

An Intro to DifferentialEquations.jl

Chris Rackauckas

September 25, 2020

0.1 Basic Introduction Via Ordinary Differential Equations

This notebook will get you started with DifferentialEquations.jl by introducing you to the functionality for solving ordinary differential equations (ODEs). The corresponding documentation page is the [ODE tutorial](#). While some of the syntax may be different for other types of equations, the same general principles hold in each case. Our goal is to give a gentle and thorough introduction that highlights these principles in a way that will help you generalize what you have learned.

0.1.1 Background

If you are new to the study of differential equations, it can be helpful to do a quick background read on [the definition of ordinary differential equations](#). We define an ordinary differential equation as an equation which describes the way that a variable u changes, that is

$$u' = f(u, p, t)$$

where p are the parameters of the model, t is the time variable, and f is the nonlinear model of how u changes. The initial value problem also includes the information about the starting value:

$$u(t_0) = u_0$$

Together, if you know the starting value and you know how the value will change with time, then you know what the value will be at any time point in the future. This is the intuitive definition of a differential equation.

0.1.2 First Model: Exponential Growth

Our first model will be the canonical exponential growth model. This model says that the rate of change is proportional to the current value, and is this:

$$u' = au$$

where we have a starting value $u(0) = u_0$. Let's say we put 1 dollar into Bitcoin which is increasing at a rate of 98% per year. Then calling now $t = 0$ and measuring time in years, our model is:

$$u' = 0.98u$$

and $u(0) = 1.0$. We encode this into Julia by noticing that, in this setup, we match the general form when

```
f(u,p,t) = 0.98u
```

```
f (generic function with 1 method)
```

with `u_0 = 1.0`. If we want to solve this model on a time span from `t=0.0` to `t=1.0`, then we define an `ODEProblem` by specifying this function `f`, this initial condition `u0`, and this time span as follows:

```
using DifferentialEquations
f(u,p,t) = 0.98u
u0 = 1.0
tspan = (0.0,1.0)
prob = ODEProblem(f,u0,tspan)
```

```
ODEProblem with uType Float64 and tType Float64. In-place: false
timespan: (0.0, 1.0)
u0: 1.0
```

To solve our `ODEProblem` we use the command `solve`.


```
sol = solve(prob)
```

```
retcode: Success
Interpolation: automatic order switching interpolation
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

and that's it: we have successfully solved our first ODE!

Analyzing the Solution Of course, the solution type is not interesting in and of itself. We want to understand the solution! The documentation page which explains in detail the functions for analyzing the solution is the [Solution Handling](#) page. Here we will describe some of the basics. You can plot the solution using the plot recipe provided by [Plots.jl](#):


```
using Plots; gr()
plot(sol)
```



From the picture we see that the solution is an exponential curve, which matches our intuition. As a plot recipe, we can annotate the result using any of the [Plots.jl attributes](#). For example:


```
plot(sol,linewidth=5,title="Solution to the linear ODE with a thick line",
      axis="Time (t)",axis="u(t) (in μm)",label="My Thick Line!") # legend=false
```

Solution to the linear ODE with a thick line



Using the mutating `plot!` command we can add other pieces to our plot. For this ODE we know that the true solution is $u(t) = u_0 \exp(at)$, so let's add some of the true solution to our plot:

```
plot!(sol.t, t->1.0*exp(0.98t),lw=3,ls=:dash,label="True Solution!")
```



In the previous command I demonstrated `sol.t`, which grabs the array of time points that the solution was saved at:

```
sol.t

5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
```

We can get the array of solution values using `sol.u`:

```
sol.u

5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

`sol.u[i]` is the value of the solution at time `sol.t[i]`. We can compute arrays of functions of the solution values using standard comprehensions, like:

```
[t+u for (u,t) in tuples(sol)]
```

```
5-element Array{Float64,1}:
 1.0
 1.2038471492789395
 1.7643769243364813
 2.66648203038311
 3.664456142481451
```

However, one interesting feature is that, by default, the solution is a continuous function. If we check the print out again:

```
sol

retcode: Success
Interpolation: automatic order switching interpolation
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

you see that it says that the solution has a order changing interpolation. The default algorithm automatically switches between methods in order to handle all types of problems. For non-stiff equations (like the one we are solving), it is a continuous function of 4th order accuracy. We can call the solution as a function of time `sol(t)`. For example, to get the value at `t=0.45`, we can use the command:

```
sol(0.45)
```

```
1 . 5 5 4 2 6 1 0 4 8 0 5 5 3 1 2
```

Controlling the Solver `DifferentialEquations.jl` has a common set of solver controls among its algorithms which can be found [at the Common Solver Options](#) page. We will detail some of the most widely used options.

The most useful options are the tolerances `abstol` and `reltol` . These tell the internal adaptive time stepping engine how precise of a solution you want. Generally, `reltol` is the relative accuracy while `abstol` is the accuracy when `u` is near zero. These tolerances are local tolerances and thus are not global guarantees. However, a good rule of thumb is that the total solution accuracy is 1-2 digits less than the relative tolerances. Thus for the defaults `abstol=1e-6` and `reltol=1e-3` , you can expect a global accuracy of about 1-2 digits. If we want to get around 6 digits of accuracy, we can use the commands:

```
sol = solve(prob, abstol=1e-8, reltol=1e-8)

retcode: Success
Interpolation: automatic order switching interpolation
t: 9-element Array{Float64,1}:
 0.0
 0.04127492324135852
 0.14679917846877366
```

```

0.28631546412766684
0.4381941361169628
0.6118924302028597
0.7985659100883337
0.9993516479536952
1.0
u: 9-element Array{Float64,1}:
1.0
1.0412786454705882
1.1547261252949712
1.3239095703537043
1.5363819257509728
1.8214895157178692
2.1871396448296223
2.662763824115295
2.664456241933517


```

Now we can see no visible difference against the true solution:

```

plot(sol)
plot!(sol.t, t->1.0*exp(0.98t),lw=3,ls=:dash,label="True Solution!")

```



Notice that by decreasing the tolerance, the number of steps the solver had to take was 9 instead of the previous 5. There is a trade off between accuracy and speed, and it is up to you to determine what is the right balance for your problem.

Another common option is to use `saveat` to make the solver save at specific time points. For example, if we want the solution at an even grid of $t=0.1k$ for integers k , we would use the command:

```

sol = solve(prob,saveat=0.1)

retcode: Success

```

```

Interpolation: 1st order linear
t: 11-element Array{Float64,1}:
 0.0
 0.1
 0.2
 0.3
 0.4
 0.5
 0.6
 0.7
 0.8
 0.9
 1.0
u: 11-element Array{Float64,1}:
 1.0
 1.102962785129292
 1.2165269512238264
 1.341783821227542
 1.4799379510586077
 1.632316207054161
 1.8003833264983584
 1.9857565541588758
 2.1902158127997695
 2.415725742084496
 2.664456142481451

```

Notice that when `saveat` is used the continuous output variables are no longer saved and thus `sol(t)`, the interpolation, is only first order. We can save at an uneven grid of points by passing a collection of values to `saveat`. For example:

```
sol = solve(prob,saveat=[0.2,0.7,0.9])
```

```

retcode: Success
Interpolation: 1st order linear
t: 3-element Array{Float64,1}:
 0.2
 0.7
 0.9
u: 3-element Array{Float64,1}:
 1.2165269512238264
 1.9857565541588758
 2.415725742084496

```

If we need to reduce the amount of saving, we can also turn off the continuous output directly via `dense=false`:

```

sol = solve(prob,dense=false)

retcode: Success
Interpolation: 1st order linear
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448

```

```
1.9730384275622996
2.664456142481451
```

and to turn off all intermediate saving we can use `save_everystep=false`:

```
sol = solve(prob,save_everystep=false)
```

```
retcode: Success
Interpolation: 1st order linear
t: 2-element Array{Float64,1}:
 0.0
 1.0
u: 2-element Array{Float64,1}:
 1.0
 2.664456142481451
```

If we want to solve and only save the final value, we can even set `save_start=false`.

```
sol = solve(prob,save_everystep=false,save_start = false)
```

```
retcode: Success
Interpolation: 1st order linear
t: 1-element Array{Float64,1}:
 1.0
u: 1-element Array{Float64,1}:
 2.664456142481451
```

Note that similarly on the other side there is `save_end=false`.

More advanced saving behaviors, such as saving functionals of the solution, are handled via the `SavingCallback` in the [Callback Library](#) which will be addressed later in the tutorial.

Choosing Solver Algorithms There is no best algorithm for numerically solving a differential equation. When you call `solve(prob)`, `DifferentialEquations.jl` makes a guess at a good algorithm for your problem, given the properties that you ask for (the tolerances, the saving information, etc.). However, in many cases you may want more direct control. A later notebook will help introduce the various *algorithms* in `DifferentialEquations.jl`, but for now let's introduce the *syntax*.

The most crucial determining factor in choosing a numerical method is the stiffness of the model. Stiffness is roughly characterized by a Jacobian `f` with large eigenvalues. That's quite mathematical, and we can think of it more intuitively: if you have big numbers in `f` (like parameters of order `1e5`), then it's probably stiff. Or, as the creator of the MATLAB ODE Suite, Lawrence Shampine, likes to define it, if the standard algorithms are slow, then it's stiff. We will go into more depth about diagnosing stiffness in a later tutorial, but for now note that if you believe your model may be stiff, you can hint this to the algorithm chooser via `alg_hints = [:stiff]`.

```
sol = solve(prob,alg_hints=[:stiff])
```

```
retcode: Success
Interpolation: specialized 3rd order "free" stiffness-aware interpolation
t: 8-element Array{Float64,1}:
 0.0
 0.05653299582822294
 0.17270731152826024
 0.3164602871490142
```



```

0.5057500163821153
0.7292241858994543
0.9912975001018789
1.0
u: 8-element Array{Float64,1}:
1.0
1.0569657840332976
1.1844199383303913
1.3636037723365293
1.6415399686182572
2.0434491434754793
2.641825616057761
2.6644526430553817

```

Stiff algorithms have to solve implicit equations and linear systems at each step so they should only be used when required.

If we want to choose an algorithm directly, you can pass the algorithm type after the problem as `solve(prob,alg)`. For example, let's solve this problem using the `Tsit5()` algorithm, and just for show let's change the relative tolerance to `1e-6` at the same time:

```

sol = solve(prob,Tsit5(),reltol=1e-6)

retcode: Success
Interpolation: specialized 4th order "free" interpolation
t: 10-element Array{Float64,1}:
0.0
0.028970819746309166
0.10049147151547619
0.19458908698515082
0.3071725081673423
0.43945421453622546
0.5883434923759523
0.7524873357619015
0.9293021330536031
1.0
u: 10-element Array{Float64,1}:
1.0
1.0287982807225062
1.1034941463604806
1.2100931078233779
1.351248605624241
1.538280340326815
1.7799346012651116
2.090571742234628
2.486102171447025
2.6644562434913377

```

0.1.3 Systems of ODEs: The Lorenz Equation

Now let's move to a system of ODEs. The [Lorenz equation](#) is the famous "butterfly attractor" that spawned chaos theory. It is defined by the system of ODEs:

$$\frac{dx}{dt} = \sigma(y - x) \quad (1)$$

$$\frac{dy}{dt} = x(\rho - z) - y \quad (2)$$

$$\frac{dz}{dt} = xy - \beta z \quad (3)$$

To define a system of differential equations in `DifferentialEquations.jl`, we define our `f` as a vector function with a vector initial condition. Thus, for the vector `u = [x,y,z]'`, we have the derivative function:

```
function lorenz!(du,u,p,t)
    σ,ρ,β = p
    du[1] = σ*(u[2]-u[1])
    du[2] = u[1]*(ρ-u[3]) - u[2]
    du[3] = u[1]*u[2] - β*u[3]
end
```

```
lorenz! (generic function with 1 method)
```

Notice here we used the in-place format which writes the output to the preallocated vector `du`. For systems of equations the in-place format is faster. We use the initial condition $u_0 = [1.0, 0.0, 0.0]$ as follows:

```
u0 = [1.0,0.0,0.0]
```

```
3-element Array{Float64,1}:
 1.0
 0.0
 0.0
```

Lastly, for this model we made use of the parameters `p`. We need to set this value in the `ODEProblem` as well. For our model we want to solve using the parameters $\sigma = 10$, $\rho = 28$, and $\beta = 8/3$, and thus we build the parameter collection:

```
p = (10,28,8/3) # we could also make this an array, or any other type!

(10, 28, 2.6666666666666665)
```

Now we generate the `ODEProblem` type. In this case, since we have parameters, we add the parameter values to the end of the constructor call. Let's solve this on a time span of `t=0` to `t=100`:

```
tspan = (0.0,100.0)
prob = ODEProblem(lorenz!,u0,tspan,p)
```

```
ODEProblem with uType Array{Float64,1} and tType Float64. In-place: true
timespan: (0.0, 100.0)
u0: [1.0, 0.0, 0.0]
```

Now, just as before, we solve the problem:

```
sol = solve(prob)

retcode: Success
Interpolation: automatic order switching interpolation
t: 1294-element Array{Float64,1}:
```

```

0.0
3.5678604836301404e-5
0.0003924646531993154
0.0032624077544510573
0.009058075635317072
0.01695646895607931
0.0276899566248403
0.041856345938267966
0.06024040228733675
0.08368539694547242

:*(99.3940307091529799.4700114749437599.5437965690901599.61465155834999.6909382314810199.787330232337
1294-element Array{Array{Float64,1},1}:
 [1.0, 0.0, 0.0]
 [0.9996434557625105, 0.0009988049817849058, 1.781434788799208e-8]
 [0.9961045497425811, 0.010965399721242457, 2.146955365838907e-6]
 [0.9693591634199452, 0.08977060667778931, 0.0001438018342266937]
 [0.9242043615038835, 0.24228912482984957, 0.0010461623302512404]
 [0.8800455868998046, 0.43873645009348244, 0.0034242593451028745]
 [0.8483309877783048, 0.69156288756671, 0.008487623500490047]
 [0.8495036595681027, 1.0145425335433382, 0.01821208597613427]
 [0.9139069079152129, 1.4425597546855036, 0.03669381053327124]
 [1.0888636764765296, 2.052326153029042, 0.07402570506414284]

:*((12.999157033749652, 14.10699925404482, 31.74244844521858) [11.646131422021162,
7.2855792145502845, 35.365000488215486] [7.777555445486692, 2.5166095828739574,
32.030953593541675] [4.739741627223412, 1.5919220588229062,
27.249779003951755] [3.2351668945618774, 2.3121727966182695,
22.724936101772805] [3.310411964698304, 4.28106626744641,
18.435441144016366] [4.527117863517627, 6.895878639772805,
16.58544600757436] [8.043672261487556, 12.71155298531689,
18.12537420595938] [9.97537965430362, 15.143884806010783, 21.00643286956427]

```

The same solution handling features apply to this case. Thus `sol.t` stores the time points and `sol.u` is an array storing the solution at the corresponding time points.

However, there are a few extra features which are good to know when dealing with systems of equations. First of all, `sol` also acts like an array. `sol[i]` returns the solution at the *i*th time point.

```

sol.t[10],sol[10]

(0.08368539694547242, [1.0888636764765296, 2.052326153029042, 0.07402570506
414284])

```

Additionally, the solution acts like a matrix where `sol[j,i]` is the value of the *j*th variable at time *i*:

```

sol[2,10]

2.052326153029042

```

We can get a real matrix by performing a conversion:

```


A = Array(sol)

3×1294 Array{Float64,2}:
 1.0  0.999643  0.996105  0.969359  ...*( 4.52712 8.04367 9.975380.0 0.000998805
0.0109654 0.0897706 6.89588 12.7116 15.14390.0 1.78143e-8 2.14696e-6 0.000143802 16.5854
18.1254 21.0064

```


This is the same as `sol`, i.e. `sol[i,j] = A[i,j]`, but now it's a true matrix. Plotting will by default show the time series for each variable:

```
plot(sol)
```




If we instead want to plot values against each other, we can use the `vars` command. Let's plot variable 1 against variable 2 against variable 3:

```
plot(sol,vars=(1,2,3))
```




This is the classic Lorenz attractor plot, where the x axis is $u[1]$, the y axis is $u[2]$, and the z axis is $u[3]$. Note that the plot recipe by default uses the interpolation, but we can turn this off:

```
plot(sol, vars=(1,2,3), denseplot=false)
```



Yikes! This shows how calculating the continuous solution has saved a lot of computational effort by computing only a sparse solution and filling in the values! Note that in vars, 0=time, and thus we can plot the time series of a single component like:

```
plot(sol, vars=(0,2))
```



0.2 Internal Types

The last basic user-interface feature to explore is the choice of types. DifferentialEquations.jl respects your input types to determine the internal types that are used. Thus since in the previous cases, when we used `Float64` values for the initial condition, this meant that the internal values would be solved using `Float64`. We made sure that time was specified via `Float64` values, meaning that time steps would utilize 64-bit floats as well. But, by simply changing these types we can change what is used internally.

As a quick example, let's say we want to solve an ODE defined by a matrix. To do this, we can simply use a matrix as input.

```
A = [1. 0 0 -5
      4 -2 4 -3
      -4 0 0 1
      5 -2 2 3]
u0 = rand(4,2)
tspan = (0.0,1.0)
f(u,p,t) = A*u
prob = ODEProblem(f,u0,tspan)
sol = solve(prob)
```

```
retcode: Success
```

```
Interpolation: automatic order switching interpolation
```

```

t: 11-element Array{Float64,1}:
 0.0
 0.04244445270868048
 0.10264882166596963
 0.16050154945121875
 0.23489487313208723
 0.33447641266750416
 0.4563000693825595
 0.6037053473032081
 0.7674242341001349
 0.9371184382013034
 1.0
u: 11-element Array{Array{Float64,2},1}:
 [0.8560802386936441 0.9041991530320039; 0.5172765489558993 0.5337473709429
 879; 0.11977614846773643 0.3282858403427533; 0.8253585894813911 0.582782007
 4680976]
 [0.6873604362006522 0.790514116484995; 0.4965866147435862 0.58602742587218
 75; 0.0282454662359211 0.2135883535498354; 1.073654349495993 0.828262138699
 8659]
 [0.3454311739890223 0.5310409175340428; 0.34559169288714264 0.536158877558
 8468; -0.023794389946542434 0.11211103915396779; 1.403452398502685 1.159774
 5839738023]
 [-0.09442415801356341 0.17371576179894965; 0.0813643033129478 0.3627856353
 1801444; 0.03462832209568199 0.10419344871562027; 1.6846702186987552 1.4495
 95296577732]
 [-0.8105893427441719 -0.4347811658386703; -0.39349422183436067 -0.01097855
 962385158; 0.3019769027914189 0.258982240513327; 1.9752244803643655 1.76238
 75912375868]
 [-1.9959667074055458 -1.4815717341582153; -1.164233162295246 -0.6885386976
 550212; 1.0628565897972229 0.8233356175871878; 2.1959838096507447 2.0350542
 958699407]
 [-3.6757234833809562 -3.020507163239456; -2.0587634945035607 -1.5632255367
 554073; 2.705477952922223 2.1653278794826045; 2.1180682303676823 2.05957618
 66650295]
 [-5.702136216614708 -4.964074842220198; -2.4984274352003046 -2.16560328043
 82965; 5.749078330092321 4.796299563435988; 1.3631889238557533 1.4913911701
 354103]
 [-7.209705068452976 -6.572419950200866; -1.2194868824181124 -1.36080072617
 97286; 10.132798777104771 8.759362371227702; -0.49333953424606247 -0.079242
 82360305845]
 [-6.77979433637902 -6.564953937575206; 3.123241849354386 2.149009759604728
 ; 14.70277867524787 13.118952350332723; -3.599222114712509 -2.8287661608196
 584]
 [-5.827606210385362 -5.870600780155404; 5.662596913011308 4.30329971685852
 3; 16.028675212109267 14.474716194196905; -5.016253403287097 -4.11082020864
 6803]

```

There is no real difference from what we did before, but now in this case `u0` is a `4x2` matrix. Because of that, the solution at each time point is matrix:

```

sol[3]

4×2 Array{Float64,2}:
 0.345431  0.531041
 0.345592  0.536159
 -0.0237944  0.112111
 1.40345    1.15977

```

In `DifferentialEquations.jl`, you can use any type that defines `+`, `-`, `*`, `/`, and has an appropriate `norm`. For example, if we want arbitrary precision floating point numbers, we can change

the input to be a matrix of BigFloat:

```
big_u0 = big.(u0)
```

```
4×@*(2 Array{*@{BigFloat,2}}:
```

```
0.85608  0.904199
0.517277 0.533747
0.119776 0.328286
0.825359 0.582782
```

and we can solve the ODEProblem with arbitrary precision numbers by using that initial condition:

```
prob = ODEProblem(f, big_u0, tspan)
sol = solve(prob)
```

```
retcode: Success
```

```
Interpolation: automatic order switching interpolation
```

```
t: 6-element Array{Float64,1}:
```

```
0.0
0.10396426099019132
0.3189803551515885
0.6045115879068522
0.926593118577873
1.0
```

```
u: 6-element Array{Array{BigFloat,2},1}:
```

```
[0.8560802386936441354947646686923690140247344970703125 0.9041991530320039
327506265181000344455242156982421875; 0.51727654895589925132526332163251936
435699462890625 0.533747370942987853226213701418600976467132568359375; 0.11
97761484677364318685022226418368518352508544921875 0.3282858403427533211527
13397168554365634918212890625; 0.825358589481391069497817625233437865972518
9208984375 0.582782007468097607016943584312684834003448486328125]
```

```
[0.33662652899518924301827783481269715939682363818748089207700124490478196
90582486 0.5240839801451442919655690779839234159403321388601296584840621174
427093175855904; 0.34080617742108863511560205610219450912309656874891003004
60607189798417586304823 0.5335237479824498628417759459587785143618427008001
843021177374669405617057343656; -0.0237381704477081309973597872310698186048
775599888962311173541258803026282910587 0.11086529762290236825827250766269
53912499740668143999669603514254684611192564207; 1.410281929256287659548493
632565838840677967196313404630277809029977510499935438 1.166721233463312303
473580419497392408405852833141444131846921967968202487494343]
```

```
[-1.7971536478788545944663846514425077063712514849170008379289436739280673
07483867 -1.303425654452695026456015360826935782500218814192194357656509737
686638383998249; -1.0403263478894363277678730592700586604331195497487962437
79237758976533464440005 -0.575215777131648509927119949283668083366922735517
6355998172791082357933499393756; 0.9114376972567753054157788622235681618320
779424507074166859520565924513198077843 0.705733945862610836097312457634447
2736878894958792950896543981157135378503667166; 2.1765191334171024624737762
04819505004292163447195961732405224276810203328516519 2.0057244272617974790
18496509006668969052934227623833401338030856075987421675913]
```

```
[-5.7122234787766846633058033857291157424391204073496066384745085208160307
052221 -4.97408602908670702980843641765136446677395844183433155090989758511
8374950871175; -2.497521221987803231698028450694716736150742317288300089298
972702271517052553165 -2.16624029344376973418040066441299378742667545336672
214203624855645215170091493; 5.76858655289688286225044543405695398334747050
16889662591578860875486932430037 4.8135305270551064821755493428733766707279
55343276835138063103738246147981874062; 1.356784573232033976131934004797674
058948880401166380984630603009924627439376629 1.486200525326082980121505651
109923300471027225312894213812095385508383325618436]
```

```
[-6.8913235332480013203924441014887759253199070595529152090574379961712500
```



```

95008493 -6.638998786932069790767753753372706360285230439984436906850507254
671483466693659; 2.74905449830727272052027393956127048093489956151836402955
6809162904097280094994 1.83572822997000909967589303241480194938094495843022
2537801112745055306921743956; 14.451685769393294340656382454395840360718142
40290820520412664889873595720692556 12.869707035719918370207129723741821221
84180485443218327295969650871953345250392; -3.37445403385254569763056254820
0304540320037390079323208400284946375854011274433 -2.6268035121341978739682
76082938739542139065502052658944200909427648303927262499]
[-5.8276015099091714221031887643341654143741285923581852581238161248801279
86560958 -5.870599492760751044235988364478636350411315078338826426479485991
659736402092067; 5.66264169951232062626029139231173855957944064491608271967
5725651341891539855325 4.30333763001209838900441736651406779916368608247840
8158162651101325158517830261; 16.028715214588039604327399636299967481999263
02038680880854096293369657654710016 14.474754325369579176486991189303984328
58126719656606013899996386712896849890289; -5.01627910748772384060686099761
5856260409024593089784259977757925578166019009524 -4.1108426719901786961413
99817401528490511017123568198587310359288145741597778994]

```

```
sol[1,3]
```

```

- 1 . 7 9 7 1 5 3 6 4 7 8 7 8 8 5 4 5 9 4 4 6 6 3 8 4 6 5 1 4 4 2 5 0 7 7 0 6 3 7 1 2 5 1 4 8 4 9 1
7 0 0 0 8 3 7 9 2 8 9 4 3 6 7 3 9 2 8 0 6 7 3 0 7 4 8 3 8 6 7

```

To really make use of this, we would want to change `abstol` and `reltol` to be small! Notice that the type for "time" is different than the type for the dependent variables, and this can be used to optimize the algorithm via keeping multiple precisions. We can convert time to be arbitrary precision as well by defining our time span with `BigFloat` variables:

```

prob = ODEProblem(f,big_u0,big.(tspan))
sol = solve(prob)

```

```

retcode: Success
Interpolation: automatic order switching interpolation
t: 6-element Array{BigFloat,1}:
 0.0
 0.103964260990191327213963623544730416089180240992277797766684309834066403
9077787
 0.318980355151588527717164445844679678027732478077720049031069688169343647
0137345
 0.604511587906852231939075257932047741356960392579240218494957170705554664
442237
 0.926593118577872971653407637786215051394119585569233450242934404857135789
1188578
 1.0
u: 6-element Array{Array{BigFloat,2},1}:
 [0.8560802386936441354947646686923690140247344970703125 0.9041991530320039
327506265181000344455242156982421875; 0.51727654895589925132526332163251936
435699462890625 0.533747370942987853226213701418600976467132568359375; 0.11
97761484677364318685022226418368518352508544921875 0.3282858403427533211527
13397168554365634918212890625; 0.825358589481391069497817625233437865972518
9208984375 0.582782007468097607016943584312684834003448486328125]
 [0.33662652899518919867759023871698550026725361316503150278813367636635034
84122464 0.5240839801451442569044419390001402995251734522331366558566501187
611431793456454; 0.34080617742108861094102647285057189307562952016031341470
89638356980030964828397 0.5335237479824498494538625894115735017580141575261
04840904336606459021439438454; -0.02373817044770813057621983918918966565464
993464902908580482374756787687290094892 0.110865297622902362119614906101173
140999345485365085227855581783856795871902727; 1.41028192925628769378661749
0508492460547425639562747366746518100418502193566071 1.16672123346331233830

```

```

8503404021018810105591350381873616351051135425652829844569]
[-1.7971536478788547846422105744845757518064754551512440564800544142703661
12312462 -1.303425654452695196418488162260366029169382372412620880399702442
946051498089127; -1.0403263478894364476315340633958498252145925019028598598
98149537362247496060962 -0.575215777131648618778063916339950762522373654615
9835408682009605229476627739613; 0.9114376972567754458777074531879734868333
925698775142155223318413774213652356813 0.705733945862610944377095325627341
233082495087184039675436864669622194094098881; 2.17651913341710248418120479
134774273383389420819689820321449470483563654539494 2.005724427261797509944
343655918737837950051953437003241294298400378307318890353]
[-5.7122234787766847958757174976551819343613448437628783356629751913525734
9708111 -4.9740860290867071614123316472574890488344202424906191533648074370
15894932962449; -2.49752122198780321949634223767974857855835874004814354635
5405241294188183759904 -2.1662402934437697423316494903242420968934504747374
55568847899987877528122674121; 5.768586552896883119045187546673165458255330
829000574276118622731700792287802841 4.813530527055106709020049316313366623
428663997007490707495037885202175278183183; 1.35678457323203389170078340391
1699455520915958270184191980968133827117916385942 1.48620052532608291167007
4554293295404186731590293709656405372580101352869516436]
[-6.8913235332480012546765871605558367374483746571689806894171407500547823
76493907 -6.638998786932069748003705987383612777597524188241769799101337096
697996974170088; 2.74905449830727295008724849978916414675801165277901591662
2831493545244870202203 1.83572822997000929148342759034406788885387917393388
6709761143339021930086297495; 14.451685769393294499932099839065922477789205
85916568436171709072435567428175799 12.869707035719918527760115937571304554
15182329995438093932840592018710982788689; -3.37445403385254583704770542596
1225423498285837178852007613195794243653839318255 -2.6268035121341979991158
23262834230743945851955340210752047557972076940585332431]
[-5.8276015099091714221032047749895607110700184867156083015459420479535862
60729707 -5.870599492760751044236003562454840273642378359938725741974453324
818158499841412; 5.66264169951232062626029098191869895623915670713783733538
3164028415485021352558 4.30333763001209838900441530056426271182709518884278
0853617423960260911780658246; 16.028715214588039604327429492339057209282919
31792323990406137557871940175730947 14.474754325369579176487017766252647897
56303802533145193195164668261023566450255; -5.01627910748772384060686634436
1327715741507466048960460303559548072307021420535 -4.1108426719901786961414
04015982016972994900647792991679621224122059843075901115]

```

Let's end by showing a more complicated use of types. For small arrays, it's usually faster to do operations on static arrays via the package [StaticArrays.jl](#). The syntax is similar to that of normal arrays, but for these special arrays we utilize the `@SMatrix` macro to indicate we want to create a static array.

```

using StaticArrays
A = @SMatrix [ 1.0  0.0 0.0 -5.0
               4.0 -2.0 4.0 -3.0
              -4.0  0.0 0.0  1.0
               5.0 -2.0 2.0  3.0]
u0 = @SMatrix rand(4,2)
tspan = (0.0,1.0)
f(u,p,t) = A*u
prob = ODEProblem(f,u0,tspan)
sol = solve(prob)

retcode: Success
Interpolation: automatic order switching interpolation
t: 10-element Array{Float64,1}:
 0.0
 0.03466471968189963

```

```

0.102924084886496
0.185727005546646
0.28613066860277514
0.4078690959069953
0.539774497821959
0.6917346336224831
0.8455349386291112
1.0
u: 10-element Array{StaticArrays.SArray{Tuple{4,2},Float64,2,8},1}:
 [0.6678548507659938 0.7442807739163952; 0.6642605720645545 0.0944976477843
4133; 0.9497717547596194 0.8870977440325551; 0.6185692703188139 0.819094165
3248551]
 [0.5650061212970365 0.6026983612782604; 0.7533652244898965 0.1973891192759
757; 0.8887291847182903 0.8261395335742894; 0.8147322009103141 1.0843475418
372746]
 [0.2534255626066652 0.1752111718163571; 0.7905278252218471 0.2231557857362
2424; 0.8418015567167034 0.8064995889158991; 1.172158854069822 1.5742338003
095222]
 [-0.3105384345849445 -0.6008125490858591; 0.6210718291275006 -0.0218867424
35127754; 0.9584509120630239 1.0214926367071773; 1.5311512895446837 2.07856
87880676074]
 [-1.233059985151846 -1.880502240847907; 0.192310959861965 -0.6083414525427
511; 1.430005152892541 1.7400718966224606; 1.810098607429582 2.491933462867
06]
 [-2.5937164002634336 -3.793827187463698; -0.4306426159030017 -1.4499236898
907477; 2.5802099124181073 3.424017625173468; 1.8369264074622103 2.58348249
8153002]
 [-4.1279910408650995 -5.9997192850577115; -0.8295592661177034 -1.980739229
7613169; 4.575878177916328 6.323038292396253; 1.3740128701237269 2.01513522
47952765]
 [-5.476712625862622 -8.037715381650798; -0.26429643874022435 -1.1774227540
103501; 7.655753601753048 10.83086674279012; 0.08811947196200043 0.30861136
18788058]
 [-5.64189953616111 -8.51189415924031; 2.176323318028735 2.2507392410318317
; 11.017722456737935 15.869629815221735; -2.077998582449022 -2.666942096674
1457]
 [-3.7021605136081677 -6.096456606288779; 7.070648130499037 9.1816742239409
29; 13.497291235530534 19.847963588894842; -5.000478936357684 -6.7890713178
21476]

sol[3]

4×@*(2 StaticArrays.SArray{@Tuple{4,2},Float64,2,8} with indices SOneTo(4)
×@*(SOneTo(2):0.253426 0.1752110.790528 0.2231560.841802 0.80651.17216
1.57423

```

0.3 Conclusion

These are the basic controls in `DifferentialEquations.jl`. All equations are defined via a problem type, and the `solve` command is used with an algorithm choice (or the default) to get a solution. Every solution acts the same, like an array `sol[i]` with `sol.t[i]`, and also like a continuous function `sol(t)` with a nice plot command `plot(sol)`. The Common Solver Options can be used to control the solver for any equation type. Lastly, the types used in the numerical solving are determined by the input types, and this can be used to solve with arbitrary precision and add additional optimizations (this can be used to solve via

GPUs for example!). While this was shown on ODEs, these techniques generalize to other types of equations as well.

0.4 Appendix

This tutorial is part of the SciMLTutorials.jl repository, found at: <https://github.com/SciML/SciMLTutorials.jl>. For more information on doing scientific machine learning (SciML) with open source software, check out <https://sciml.ai/>.

To locally run this tutorial, do the following commands:

```
using SciMLTutorials
SciMLTutorials.weave_file("introduction", "01-ode_introduction.jmd")
```

Computer Information:

```
Julia Version 1.4.2
Commit 44fa15b150* (2020-05-23 18:35 UTC)
Platform Info:
  OS: Linux (x86_64-pc-linux-gnu)
  CPU: Intel(R) Core(TM) i7-9700K CPU @ 3.60GHz
  WORD_SIZE: 64
  LIBM: libopenlibm
  LLVM: libLLVM-8.0.1 (ORCJIT, skylake)
```

Environment:

```
JULIA_LOAD_PATH = /builds/JuliaGPU/DiffEqTutorials.jl:
JULIA_DEPOT_PATH = /builds/JuliaGPU/DiffEqTutorials.jl/.julia
JULIA_CUDA_MEMORY_LIMIT = 2147483648
JULIA_NUM_THREADS = 8
```

Package Information:

```
Status `~/builds/JuliaGPU/DiffEqTutorials.jl/tutorials/introduction/Project.toml`
[6e4b80f9-dd63-53aa-95a3-0cdb28fa8baf] BenchmarkTools 0.5.0
[0c46a032-eb83-5123-abaf-570d42b7fbaf] DifferentialEquations 6.15.0
[65888b18-ceab-5e60-b2b9-181511a3b968] ParameterizedFunctions 5.6.0
[91a5bcdd-55d7-5caf-9e0b-520d859cae80] Plots 1.6.6
[90137ffa-7385-5640-81b9-e52037218182] StaticArrays 0.12.4
[c3572dad-4567-51f8-b174-8c6c989267f4] Sundials 4.3.0
[37e2e46d-f89d-539d-b4ee-838fcccc9c8e] LinearAlgebra
```