

# An Intro to DifferentialEquations.jl

Chris Rackauckas

September 13, 2020

## 0.1 Basic Introduction Via Ordinary Differential Equations

This notebook will get you started with DifferentialEquations.jl by introducing you to the functionality for solving ordinary differential equations (ODEs). The corresponding documentation page is the [ODE tutorial](#). While some of the syntax may be different for other types of equations, the same general principles hold in each case. Our goal is to give a gentle and thorough introduction that highlights these principles in a way that will help you generalize what you have learned.

### 0.1.1 Background

If you are new to the study of differential equations, it can be helpful to do a quick background read on [the definition of ordinary differential equations](#). We define an ordinary differential equation as an equation which describes the way that a variable  $u$  changes, that is

$$u' = f(u, p, t)$$

where  $p$  are the parameters of the model,  $t$  is the time variable, and  $f$  is the nonlinear model of how  $u$  changes. The initial value problem also includes the information about the starting value:

$$u(t_0) = u_0$$

Together, if you know the starting value and you know how the value will change with time, then you know what the value will be at any time point in the future. This is the intuitive definition of a differential equation.

### 0.1.2 First Model: Exponential Growth

Our first model will be the canonical exponential growth model. This model says that the rate of change is proportional to the current value, and is this:

$$u' = au$$

where we have a starting value  $u(0) = u_0$ . Let's say we put 1 dollar into Bitcoin which is increasing at a rate of 98% per year. Then calling now  $t = 0$  and measuring time in years, our model is:

$$u' = 0.98u$$

and  $u(0) = 1.0$ . We encode this into Julia by noticing that, in this setup, we match the general form when

```
f(u,p,t) = 0.98u
```

```
f (generic function with 1 method)
```

with `u_0 = 1.0`. If we want to solve this model on a time span from `t=0.0` to `t=1.0`, then we define an `ODEProblem` by specifying this function `f`, this initial condition `u0`, and this time span as follows:

```
using DifferentialEquations
f(u,p,t) = 0.98u
u0 = 1.0
tspan = (0.0,1.0)
prob = ODEProblem(f,u0,tspan)
```

```
ODEProblem with uType Float64 and tType Float64. In-place: false
timespan: (0.0, 1.0)
u0: 1.0
```

To solve our `ODEProblem` we use the command `solve`.


```
sol = solve(prob)
```

```
retcode: Success
Interpolation: automatic order switching interpolation
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

and that's it: we have successfully solved our first ODE!

**Analyzing the Solution** Of course, the solution type is not interesting in and of itself. We want to understand the solution! The documentation page which explains in detail the functions for analyzing the solution is the [Solution Handling](#) page. Here we will describe some of the basics. You can plot the solution using the plot recipe provided by [Plots.jl](#):


```
using Plots; gr()
plot(sol)
```



From the picture we see that the solution is an exponential curve, which matches our intuition. As a plot recipe, we can annotate the result using any of the [Plots.jl attributes](#). For example:

```
plot(sol,linewidth=5,title="Solution to the linear ODE with a thick line",
      axis="Time (t)",axis="u(t) (in  $\mu\text{m}$ )",label="My Thick Line!") # legend=false
```


### Solution to the linear ODE with a thick line



Using the mutating `plot!` command we can add other pieces to our plot. For this ODE we know that the true solution is  $u(t) = u_0 \exp(at)$ , so let's add some of the true solution to our plot:

```
plot!(sol.t, t->1.0*exp(0.98t),lw=3,ls=:dash,label="True Solution!")
```

## Solution to the linear ODE with a thick line



In the previous command I demonstrated `sol.t`, which grabs the array of time points that the solution was saved at:

```
sol.t
```

```
5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
```

We can get the array of solution values using `sol.u`:

```
sol.u
```

```
5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

`sol.u[i]` is the value of the solution at time `sol.t[i]`. We can compute arrays of functions of the solution values using standard comprehensions, like:

```
[t+u for (u,t) in tuples(sol)]
```

```
5-element Array{Float64,1}:
 1.0
 1.2038471492789395
 1.7643769243364813
 2.66648203038311
 3.664456142481451
```

However, one interesting feature is that, by default, the solution is a continuous function. If we check the print out again:

```
sol

retcode: Success
Interpolation: automatic order switching interpolation
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

you see that it says that the solution has a order changing interpolation. The default algorithm automatically switches between methods in order to handle all types of problems. For non-stiff equations (like the one we are solving), it is a continuous function of 4th order accuracy. We can call the solution as a function of time `sol(t)`. For example, to get the value at `t=0.45`, we can use the command:

```
sol(0.45)
```

```
1 . 5 5 4 2 6 1 0 4 8 0 5 5 3 1 2
```

**Controlling the Solver** `DifferentialEquations.jl` has a common set of solver controls among its algorithms which can be found [at the Common Solver Options](#) page. We will detail some of the most widely used options.

The most useful options are the tolerances  `abstol`  and  `reltol` . These tell the internal adaptive time stepping engine how precise of a solution you want. Generally,  `reltol`  is the relative accuracy while  `abstol`  is the accuracy when  `u`  is near zero. These tolerances are local tolerances and thus are not global guarantees. However, a good rule of thumb is that the total solution accuracy is 1-2 digits less than the relative tolerances. Thus for the defaults  `abstol=1e-6`  and  `reltol=1e-3` , you can expect a global accuracy of about 1-2 digits. If we want to get around 6 digits of accuracy, we can use the commands:

```
sol = solve(prob, abstol=1e-8, reltol=1e-8)

retcode: Success
Interpolation: automatic order switching interpolation
t: 9-element Array{Float64,1}:
 0.0
 0.04127492324135852
 0.14679917846877366
```

```

0.28631546412766684
0.4381941361169628
0.6118924302028597
0.7985659100883337
0.9993516479536952
1.0
u: 9-element Array{Float64,1}:
1.0
1.0412786454705882
1.1547261252949712
1.3239095703537043
1.5363819257509728
1.8214895157178692
2.1871396448296223
2.662763824115295
2.664456241933517


```

Now we can see no visible difference against the true solution:

```

plot(sol)
plot!(sol.t, t->1.0*exp(0.98t),lw=3,ls=:dash,label="True Solution!")

```



Notice that by decreasing the tolerance, the number of steps the solver had to take was 9 instead of the previous 5. There is a trade off between accuracy and speed, and it is up to you to determine what is the right balance for your problem.

Another common option is to use `saveat` to make the solver save at specific time points. For example, if we want the solution at an even grid of  $t=0.1k$  for integers  $k$ , we would use the command:

```

sol = solve(prob,saveat=0.1)

retcode: Success

```

```

Interpolation: 1st order linear
t: 11-element Array{Float64,1}:
 0.0
 0.1
 0.2
 0.3
 0.4
 0.5
 0.6
 0.7
 0.8
 0.9
 1.0
u: 11-element Array{Float64,1}:
 1.0
 1.102962785129292
 1.2165269512238264
 1.341783821227542
 1.4799379510586077
 1.632316207054161
 1.8003833264983584
 1.9857565541588758
 2.1902158127997695
 2.415725742084496
 2.664456142481451

```

Notice that when `saveat` is used the continuous output variables are no longer saved and thus `sol(t)`, the interpolation, is only first order. We can save at an uneven grid of points by passing a collection of values to `saveat`. For example:

```
sol = solve(prob,saveat=[0.2,0.7,0.9])
```

```

retcode: Success
Interpolation: 1st order linear
t: 3-element Array{Float64,1}:
 0.2
 0.7
 0.9
u: 3-element Array{Float64,1}:
 1.2165269512238264
 1.9857565541588758
 2.415725742084496

```

If we need to reduce the amount of saving, we can also turn off the continuous output directly via `dense=false`:

```

sol = solve(prob,dense=false)

retcode: Success
Interpolation: 1st order linear
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448

```

```
1.9730384275622996
2.664456142481451
```

and to turn off all intermediate saving we can use `save_everystep=false`:

```
sol = solve(prob,save_everystep=false)
```

```
retcode: Success
Interpolation: 1st order linear
t: 2-element Array{Float64,1}:
 0.0
 1.0
u: 2-element Array{Float64,1}:
 1.0
 2.664456142481451
```

If we want to solve and only save the final value, we can even set `save_start=false`.

```
sol = solve(prob,save_everystep=false,save_start = false)
```

```
retcode: Success
Interpolation: 1st order linear
t: 1-element Array{Float64,1}:
 1.0
u: 1-element Array{Float64,1}:
 2.664456142481451
```

Note that similarly on the other side there is `save_end=false`.

More advanced saving behaviors, such as saving functionals of the solution, are handled via the `SavingCallback` in the [Callback Library](#) which will be addressed later in the tutorial.

**Choosing Solver Algorithms** There is no best algorithm for numerically solving a differential equation. When you call `solve(prob)`, `DifferentialEquations.jl` makes a guess at a good algorithm for your problem, given the properties that you ask for (the tolerances, the saving information, etc.). However, in many cases you may want more direct control. A later notebook will help introduce the various *algorithms* in `DifferentialEquations.jl`, but for now let's introduce the *syntax*.

The most crucial determining factor in choosing a numerical method is the stiffness of the model. Stiffness is roughly characterized by a Jacobian `f` with large eigenvalues. That's quite mathematical, and we can think of it more intuitively: if you have big numbers in `f` (like parameters of order `1e5`), then it's probably stiff. Or, as the creator of the MATLAB ODE Suite, Lawrence Shampine, likes to define it, if the standard algorithms are slow, then it's stiff. We will go into more depth about diagnosing stiffness in a later tutorial, but for now note that if you believe your model may be stiff, you can hint this to the algorithm chooser via `alg_hints = [:stiff]`.

```
sol = solve(prob,alg_hints=[:stiff])
```

```
retcode: Success
Interpolation: specialized 3rd order "free" stiffness-aware interpolation
t: 8-element Array{Float64,1}:
 0.0
 0.05653299582822294
 0.17270731152826024
 0.3164602871490142
```



```

0.5057500163821153
0.7292241858994543
0.9912975001018789
1.0
u: 8-element Array{Float64,1}:
1.0
1.0569657840332976
1.1844199383303913
1.3636037723365293
1.6415399686182572
2.0434491434754793
2.641825616057761
2.6644526430553817

```

Stiff algorithms have to solve implicit equations and linear systems at each step so they should only be used when required.

If we want to choose an algorithm directly, you can pass the algorithm type after the problem as `solve(prob,alg)`. For example, let's solve this problem using the `Tsit5()` algorithm, and just for show let's change the relative tolerance to `1e-6` at the same time:

```

sol = solve(prob,Tsit5(),reltol=1e-6)

retcode: Success
Interpolation: specialized 4th order "free" interpolation
t: 10-element Array{Float64,1}:
0.0
0.028970819746309166
0.10049147151547619
0.19458908698515082
0.3071725081673423
0.43945421453622546
0.5883434923759523
0.7524873357619015
0.9293021330536031
1.0
u: 10-element Array{Float64,1}:
1.0
1.0287982807225062
1.1034941463604806
1.2100931078233779
1.351248605624241
1.538280340326815
1.7799346012651116
2.090571742234628
2.486102171447025
2.6644562434913377

```

### 0.1.3 Systems of ODEs: The Lorenz Equation

Now let's move to a system of ODEs. The [Lorenz equation](#) is the famous "butterfly attractor" that spawned chaos theory. It is defined by the system of ODEs:

$$\frac{dx}{dt} = \sigma(y - x) \quad (1)$$

$$\frac{dy}{dt} = x(\rho - z) - y \quad (2)$$

$$\frac{dz}{dt} = xy - \beta z \quad (3)$$

To define a system of differential equations in `DifferentialEquations.jl`, we define our `f` as a vector function with a vector initial condition. Thus, for the vector `u = [x,y,z]'`, we have the derivative function:

```
function lorenz!(du,u,p,t)
    σ,ρ,β = p
    du[1] = σ*(u[2]-u[1])
    du[2] = u[1]*(ρ-u[3]) - u[2]
    du[3] = u[1]*u[2] - β*u[3]
end
```

```
lorenz! (generic function with 1 method)
```

Notice here we used the in-place format which writes the output to the preallocated vector `du`. For systems of equations the in-place format is faster. We use the initial condition  $u_0 = [1.0, 0.0, 0.0]$  as follows:

```
u0 = [1.0,0.0,0.0]
```

```
3-element Array{Float64,1}:
 1.0
 0.0
 0.0
```

Lastly, for this model we made use of the parameters `p`. We need to set this value in the `ODEProblem` as well. For our model we want to solve using the parameters  $\sigma = 10$ ,  $\rho = 28$ , and  $\beta = 8/3$ , and thus we build the parameter collection:

```
p = (10,28,8/3) # we could also make this an array, or any other type!

(10, 28, 2.6666666666666665)
```

Now we generate the `ODEProblem` type. In this case, since we have parameters, we add the parameter values to the end of the constructor call. Let's solve this on a time span of `t=0` to `t=100`:

```
tspan = (0.0,100.0)
prob = ODEProblem(lorenz!,u0,tspan,p)
```

```
ODEProblem with uType Array{Float64,1} and tType Float64. In-place: true
timespan: (0.0, 100.0)
u0: [1.0, 0.0, 0.0]
```

Now, just as before, we solve the problem:

```
sol = solve(prob)

retcode: Success
Interpolation: automatic order switching interpolation
t: 1294-element Array{Float64,1}:
```

```

0.0
3.5678604836301404e-5
0.0003924646531993154
0.0032624077544510573
0.009058075635317072
0.01695646895607931
0.0276899566248403
0.041856345938267966
0.06024040228733675
0.08368539694547242

:*(99.3940307091529799.4700114749437599.5437965690901599.61465155834999.6909382314810199.787330232337
1294-element Array{Array{Float64,1},1}:
 [1.0, 0.0, 0.0]
 [0.9996434557625105, 0.0009988049817849058, 1.781434788799208e-8]
 [0.9961045497425811, 0.010965399721242457, 2.146955365838907e-6]
 [0.9693591634199452, 0.08977060667778931, 0.0001438018342266937]
 [0.9242043615038835, 0.24228912482984957, 0.0010461623302512404]
 [0.8800455868998046, 0.43873645009348244, 0.0034242593451028745]
 [0.8483309877783048, 0.69156288756671, 0.008487623500490047]
 [0.8495036595681027, 1.0145425335433382, 0.01821208597613427]
 [0.9139069079152129, 1.4425597546855036, 0.03669381053327124]
 [1.0888636764765296, 2.052326153029042, 0.07402570506414284]

:*( [12.999157033749652, 14.10699925404482, 31.74244844521858] [11.646131422021162,
7.2855792145502845, 35.365000488215486] [7.777555445486692, 2.5166095828739574,
32.030953593541675] [4.739741627223412, 1.5919220588229062,
27.249779003951755] [3.2351668945618774, 2.3121727966182695,
22.724936101772805] [3.310411964698304, 4.28106626744641,
18.435441144016366] [4.527117863517627, 6.895878639772805,
16.58544600757436] [8.043672261487556, 12.711555298531689,
18.12537420595938] [9.97537965430362, 15.143884806010783, 21.00643286956427]

```

The same solution handling features apply to this case. Thus `sol.t` stores the time points and `sol.u` is an array storing the solution at the corresponding time points.

However, there are a few extra features which are good to know when dealing with systems of equations. First of all, `sol` also acts like an array. `sol[i]` returns the solution at the *i*th time point.

```

sol.t[10],sol[10]

(0.08368539694547242, [1.0888636764765296, 2.052326153029042, 0.07402570506
414284])

```

Additionally, the solution acts like a matrix where `sol[j,i]` is the value of the *j*th variable at time *i*:

```

sol[2,10]

2.052326153029042

```

We can get a real matrix by performing a conversion:

```


A = Array(sol)

3×1294 Array{Float64,2}:
 1.0  0.999643  0.996105  0.969359  ...*( 4.52712 8.04367 9.975380.0 0.000998805
0.0109654 0.0897706 6.89588 12.7116 15.14390.0 1.78143e-8 2.14696e-6 0.000143802 16.5854
18.1254 21.0064

```


This is the same as `sol`, i.e. `sol[i,j] = A[i,j]`, but now it's a true matrix. Plotting will by default show the time series for each variable:

```
plot(sol)
```




If we instead want to plot values against each other, we can use the `vars` command. Let's plot variable 1 against variable 2 against variable 3:

```
plot(sol,vars=(1,2,3))
```




This is the classic Lorenz attractor plot, where the x axis is  $u[1]$ , the y axis is  $u[2]$ , and the z axis is  $u[3]$ . Note that the plot recipe by default uses the interpolation, but we can turn this off:

```
plot(sol, vars=(1,2,3), denseplot=false)
```



Yikes! This shows how calculating the continuous solution has saved a lot of computational effort by computing only a sparse solution and filling in the values! Note that in vars, 0=time, and thus we can plot the time series of a single component like:

```
plot(sol, vars=(0,2))
```



## 0.2 Internal Types

The last basic user-interface feature to explore is the choice of types. DifferentialEquations.jl respects your input types to determine the internal types that are used. Thus since in the previous cases, when we used `Float64` values for the initial condition, this meant that the internal values would be solved using `Float64`. We made sure that time was specified via `Float64` values, meaning that time steps would utilize 64-bit floats as well. But, by simply changing these types we can change what is used internally.

As a quick example, let's say we want to solve an ODE defined by a matrix. To do this, we can simply use a matrix as input.

```
A = [1. 0 0 -5
      4 -2 4 -3
      -4 0 0 1
      5 -2 2 3]
u0 = rand(4,2)
tspan = (0.0,1.0)
f(u,p,t) = A*u
prob = ODEProblem(f,u0,tspan)
sol = solve(prob)
```

```
retcode: Success
Interpolation: automatic order switching interpolation
```

```

t: 11-element Array{Float64,1}:
 0.0
 0.0378804853136824
 0.09822979107923621
 0.17164916169410643
 0.26696228936116195
 0.3756626403200669
 0.48679521575542184
 0.6235274643467611
 0.7784864259417906
 0.939666577559981
 1.0
u: 11-element Array{Array{Float64,2},1}:
 [0.6107426767545865 0.20450179835791737; 0.9719930975314253 0.319576889858
 14614; 0.4864414730703177 0.07577706275657703; 0.02899231113673495 0.603935
 0807702066]
 [0.6208818206442575 0.08752175289132093; 1.047676831635972 0.2577817209687
 7955; 0.39558935269469003 0.077908489214442; 0.11053170656595851 0.68913765
 61025813]
 [0.6056010176824679 -0.1401512734970753; 1.1125271414428957 0.117186890798
 1449; 0.25726939056803333 0.12853506181806693; 0.23628219333285186 0.806968
 3107562626]
 [0.5345817241725646 -0.47958716768074816; 1.1029365655916277 -0.1062917915
 149994; 0.11108883944082443 0.28138882838434215; 0.37810707051410214 0.9128
 409817855674]
 [0.35893028706650737 -1.0025869403977175; 0.9610136034112924 -0.4353033718
 5806; -0.01835272155050849 0.6520442580731975; 0.5347074395748432 0.9716898
 274832753]
 [0.0546588423246282 -1.6637222649438865; 0.6602200461256789 -0.76926162668
 97222; -0.04641085646533329 1.3337566336822007; 0.6602028273008376 0.902196
 8373556304]
 [-0.34602173129195496 -2.3260767004378144; 0.2666705000678275 -0.928815558
 6079767; 0.09266590615386605 2.3110393163001586; 0.7116170861914148 0.64825
 93647391249]
 [-0.9027100174961151 -2.9458359577621116; -0.2066110587128041 -0.644258257
 2639444; 0.5271748560738192 3.818791600153346; 0.6412531012408232 0.0431329
 169338428]
 [-1.4890750914796622 -3.057857313386042; -0.5071649623352912 0.64084188157
 46995; 1.3552354348497233 5.649646193347419; 0.34940867159267774 -1.0540331
 651213428]
 [-1.8336820932915607 -2.038430639883; -0.2746778555074971 3.34065236004437
 5; 2.4621650821609875 7.078446010371258; -0.2172199667903263 -2.59849798840
 232]
 [-1.837992657909004 -1.2582683653031514; 0.007609609620422231 4.7221368083
 04371; 2.885532656070101 7.304951139776333; -0.4971719330606758 -3.24476164
 52632357]

```

There is no real difference from what we did before, but now in this case `u0` is a `4x2` matrix. Because of that, the solution at each time point is matrix:

```
sol[3]
```

```

4×2 Array{Float64,2}:
 0.605601  -0.140151
 1.11253   0.117187
 0.257269  0.128535
 0.236282  0.806968

```

In `DifferentialEquations.jl`, you can use any type that defines `+`, `-`, `*`, `/`, and has an appropriate `norm`. For example, if we want arbitrary precision floating point numbers, we can change

the input to be a matrix of BigFloat:

```
big_u0 = big.(u0)
```

```
4×@*(2 Array{*@{BigFloat,2}}:
```

```
0.610743  0.204502
0.971993  0.319577
0.486441  0.0757771
0.0289923 0.603935
```

and we can solve the ODEProblem with arbitrary precision numbers by using that initial condition:

```
prob = ODEProblem(f, big_u0, tspan)
sol = solve(prob)
```

```
retcode: Success
```

```
Interpolation: automatic order switching interpolation
```

```
t: 6-element Array{Float64,1}:
```

```
0.0
0.03788048531368239
0.20467587727168945
0.4744527298129571
0.7704521341216637
1.0
```

```
u: 6-element Array{Array{BigFloat,2},1}:
```

```
[0.6107426767545864887409834409481845796108245849609375 0.2045017983579173
748154289569356478750705718994140625; 0.97199309753142526524527511355699971
31824493408203125 0.3195768898581461403551884359330870211124420166015625; 0
.486441473070317709215260038035921752452850341796875 0.07577706275657702938
98594900383614003658294677734375; 0.028992311136734949528204197122249752283
0963134765625 0.6039350807702066337157020825543440878391265869140625]
```

```
[0.62088182065828320350446127382857097510162188905271906732125148527226351
5770513 0.08752175286248419368139883325417646899018283960641124565048259116
681480416185191; 1.04767683176452708390684510325030658165282306239559202169
9260673556972690421968 0.25778172111354580226926353467701939843111902080601
94571160056670925065760168003; 0.395589352722751562652165617891223614264428
8119216503505248114846964958293227088 0.07790848933205268441061885995589786
555585833165028508438664432353445126043575794; 0.11053170659893504749525769
05252997936470807616135619751222795594842331310187088 0.6891376561012090218
937177211709570362794258135609415254587120977856919954588972]
```

```
[0.48412755906926765830933764199430569613547423669595562627362419734029484
09768495 -0.651728525209022002911439699661894044379953352330375323043627829
410055871702399; 1.06915058031544796926793310181482800940365937515899525072
1271580547821279994499 -0.2184262312868074514882283873453102485172308494871
567462198313012485411059732878; 0.05713671390848955356798270455000876170987
987214210650081731703803159622680639521 0.386694187641785030261076437253624
3626964642427673716620870570063732093679614671; 0.4364349416835874148824096
665527338843175754981171597827263334193629939596635352 0.944308389516635576
5938377146938971292157274917719578448001675279593002074797444]
```

```
[-0.2981652608089523266643266244599309881813704861884858326450976474256052
628534215 -2.25659500299521162949410927189812877908981656701307227404467249
229761126220399; 0.31208608754183446371020501762629904961340244366492194620
56624599842155339645621 -0.924942309409075863496899536862360074849461888175
9550328269873953942506054269207; 0.0679906893216031100744015486084956669310
5261986348070279456486381864511832986715 2.18967068707936834016744565988132
8124359894281156154943982496411704953986880742; 0.7103271818406303404572858
936830896505370058059045079889688721207225539274997665 0.686518849837398518
3975045624260579872712595858820405363289676796060105818968788]
```

```
[-1.4627553295406209024148156290005506228731406239158859431435260657006783
```



```

67911558 -3.074210909026630096493735459469761388450863179852393457045206000
838088236535704; -0.5019053856402683043709078222695444053090074162099882726
753842843956657577686607 0.544196498877584135835489822275808133166857090939
6124825609681599394665867021476; 1.3049118562571530355299090644344936334085
13710165526822336385476139406638459792 5.5593088499403390744614464820120786
29509978828332306368018275053538411113091068; 0.370541746155783269884749280
8460741848242022939780504467103399494364631874069913 -0.9868079894393118506
406714891311690561774127911195607511458531537678770199141711]
[-1.8379941447936739473612537234832771096625328266808858852830524592161109
61807102 -1.258261205079914864606808516015105125197862336475965917047756447
103748396369034; 0.00761319052454350422140926678484116959965555693946944764
9227235805828015946480803 4.72215572140100767489971838163100503248764706762
1269699979901871243727659504695; 2.8855389055355075348560060978800383396937
58466736340087304030903039705823837946 7.3049591784413509222077082218848881
25352265928420269505515636056252473239718279; -0.49717494381318145874734641
51805657631380839232493789670712244186043812457472982 -3.244771031621865588
683350107440929960194231278388141958263106996057985471395179]

```

```
sol[1,3]
```

```

0 . 4 8 4 1 2 7 5 5 9 0 6 9 2 6 7 6 5 8 3 0 9 3 3 7 6 4 1 9 9 4 3 0 5 6 9 6 1 3 5 4 7 4 2 3 6 6 9
5 9 5 5 6 2 6 2 7 3 6 2 4 1 9 7 3 4 0 2 9 4 8 4 0 9 7 6 8 4 9 5

```

To really make use of this, we would want to change `abstol` and `reltol` to be small! Notice that the type for "time" is different than the type for the dependent variables, and this can be used to optimize the algorithm via keeping multiple precisions. We can convert time to be arbitrary precision as well by defining our time span with `BigFloat` variables:

```

prob = ODEProblem(f,big_u0,big.(tspan))
sol = solve(prob)

```

```

retcode: Success
Interpolation: automatic order switching interpolation
t: 6-element Array{BigFloat,1}:
 0.0
 0.037880485313682388398084516187830238264143717656443217945847242874606236
07108105
 0.204675877271689452417522856815077613517093133111538903522066929440925188
0275945
 0.474452729812957085027989974099192943137032342325685064163471482391289143
9940861
 0.770452134121663686554807451511767273978335027743544281845128143757924344
2405248
 1.0
u: 6-element Array{Array{BigFloat,2},1}:
 [0.6107426767545864887409834409481845796108245849609375 0.2045017983579173
748154289569356478750705718994140625; 0.97199309753142526524527511355699971
31824493408203125 0.3195768898581461403551884359330870211124420166015625; 0
.486441473070317709215260038035921752452850341796875 0.07577706275657702938
98594900383614003658294677734375; 0.028992311136734949528204197122249752283
0963134765625 0.6039350807702066337157020825543440878391265869140625]
 [0.62088182065828320332529651705740607235058264120157818171302869861830593
97874694 0.0875217528624842025004562215316929451125748898883193632774130789
2259811848597719; 1.0476768317645270796027494817952786802232167515861540998
75109976298662978534304 0.2577817211135458073147732968881930768715357034600
330613925289708786167832336579; 0.39558935272275156888401385856177319378858
78058416098154106488275093265236792879 0.0779084893320526835202201652818041
127885767306264940384327995657555167288700942; 0.11053170659893504189674895
80598460248586820984894772255158577098406586052761087 0.6891376561012090162

```

```

598977316937836797517307124275254689136207606613472100123185]
[0.48412755906926768587374651732698916689769812493841005877781189322829321
88191398 -0.651728525209021915687110054830072292866588002329803625140349494
5420398036534489; 1.0691505803154479900875564615119510989344695319319249729
54877843339619302088322 -0.218426231286807395383551222455320424010603924166
750917018363777036824656300138; 0.05713671390848957791871335355671691152520
482633486200135327755241502295513109082 0.386694187641784972614061877971458
8499314242927742398453563315235512328853572722; 0.4364349416835873871902837
508366362744180323000828370421299925286838732974521927 0.944308389516635563
8586652956068834177918834355068680406961317050965641934718464]
[-0.2981652608089521847491391568000536606791273058963481243749172746687122
134935859 -2.25659500299521141977357649441640470341148710187190385315327132
20655263332742; 0.312086087541834599213104829033597991572589950945003104633
3818182314590438578674 -0.9249423094090758458996467082590674750519757337178
030205269609212397571698676065; 0.06799068932160303992456243666258417742695
2318311945301535155318681397458686753 2.18967068707936798212093799583238322
2778249497840809479348734347083860792334767; 0.7103271818406303348554222237
856463557966580523895435305277285784773645909643107 0.686518849837398628772
397458447388712054507452924717703555507082536710490026214]
[-1.4627553295406206615152431623230262169708817866732718096202535306202895
10747762 -3.074210909026630231627796525666063646569821738923088045715331302
359026109275258; -0.5019053856402682506619463876418672887796649965167761432
877240899615869819966747 0.544196498877583277552096215333386864875118321127
4352698613881570972356026300148; 1.3049118562571525834749659220698956696404
60551162890070108725423520729391309072 5.5593088499403382526812368277594478
27161520358773036570478164837703601931914988; 0.370541746155783457964801396
1560326645976878249540980796548696558136062094042815 -0.9868079894393112474
766695000241290129115073494498562235523266979337060978720073]
[-1.8379941447936739923170128812322815223914777290887843729589544693911087
81868799 -1.258261205079915903053513388654487339735294153543247920343471905
696749398837918; 0.00761319052454311103054356011730151480677218700674602423
3790435634479505521652825 4.72215572140100597648100215773131298568109658617
7140838608748871182612919061927; 2.8855389055355070592086145355497611991255
064377392159344203543755361049183776 7.304959178441350798121343395106225935
638367179305808804950156319776818775972458; -0.4971749438131811169621287074
142322640249736530973214447699170414300469436768434 -3.24477103162186483511
8642248678212050262336179988448580041235161185398298913051]

```

Let's end by showing a more complicated use of types. For small arrays, it's usually faster to do operations on static arrays via the package [StaticArrays.jl](#). The syntax is similar to that of normal arrays, but for these special arrays we utilize the `@SMatrix` macro to indicate we want to create a static array.

```

using StaticArrays
A = @SMatrix [ 1.0  0.0 0.0 -5.0
               4.0 -2.0 4.0 -3.0
              -4.0  0.0 0.0  1.0
               5.0 -2.0 2.0  3.0]
u0 = @SMatrix rand(4,2)
tspan = (0.0,1.0)
f(u,p,t) = A*u
prob = ODEProblem(f,u0,tspan)
sol = solve(prob)

retcode: Success
Interpolation: automatic order switching interpolation
t: 10-element Array{Float64,1}:
 0.0
 0.04439183904555069

```

```

0.11771676651298536
0.20605533246155705
0.31468451138166686
0.43028172071784404
0.5494312304478333
0.672171629988903
0.8246118445270447
1.0
u: 10-element Array{StaticArrays.SArray{Tuple{4,2},Float64,2,8},1}:
 [0.22932371025126863 0.1725729837929979; 0.5988783719673871 0.558270700518
6731; 0.7950484972703886 0.7067286061753983; 0.9088468667880318 0.553382553
293603]
 [0.012849075067361881 0.04120030776233033; 0.5776993047470756 0.5713125721
478617; 0.8172217810810329 0.7145206182431396; 1.0884374943986317 0.6718917
45070644]
 [-0.4485344285794179 -0.24299277221050583; 0.4315666049040173 0.5159508961
853313; 0.967321781916795 0.7976178546428485; 1.3335115454246012 0.83264558
26307495]
 [-1.152993998379832 -0.6805927715048965; 0.12683058328874391 0.35805731722
90347; 1.3729898036383266 1.0375943437650759; 1.5180056218535884 0.95192274
93592948]
 [-2.1723218101647292 -1.3143736547261546; -0.3025827139347661 0.1173815341
6315672; 2.2591195942987046 1.5736374802711457; 1.5302998828056433 0.956179
4039638346]
 [-3.301824520091458 -2.0113103581626004; -0.5918001261556982 -0.0430202146
9409953; 3.6890562112843694 2.444858603592143; 1.2212279017272847 0.7496526
190190814]
 [-4.290514628030032 -2.6074964964195826; -0.399819587666589 0.094254962675
58744; 5.6173876382208645 3.6176975282854302; 0.496547190936951 0.274570300
4772448]
 [-4.813542437416433 -2.89270731213996; 0.6592220748244556 0.77406848603044
05; 7.8714897040049046 4.9747580993720195; -0.7145782579071296 -0.509711845
4606824]
 [-4.207323203709997 -2.4300106541360664; 3.628222244181996 2.6504360523439
96; 10.445825337370675 6.477356765714792; -2.8433862820153646 -1.8701671582
185273]
 [-0.9051234009059521 -0.23176126746763082; 9.605171013153203 6.38692509666
0997; 11.66960663241951 7.03927248529625; -5.881756735513337 -3.77678947507
48603]

sol[3]

4×@*(2 StaticArrays.SArray{Tuple{4,2},Float64,2,8} with indices SOneTo(4)
×@*(SOneTo(2):-0.448534 -0.2429930.431567 0.5159510.967322 0.7976181.33351
0.832646

```

### 0.3 Conclusion

These are the basic controls in DifferentialEquations.jl. All equations are defined via a problem type, and the `solve` command is used with an algorithm choice (or the default) to get a solution. Every solution acts the same, like an array `sol[i]` with `sol.t[i]`, and also like a continuous function `sol(t)` with a nice plot command `plot(sol)`. The Common Solver Options can be used to control the solver for any equation type. Lastly, the types used in the numerical solving are determined by the input types, and this can be used to solve with arbitrary precision and add additional optimizations (this can be used to solve via

GPUs for example!). While this was shown on ODEs, these techniques generalize to other types of equations as well.

## 0.4 Appendix

This tutorial is part of the SciMLTutorials.jl repository, found at: <https://github.com/SciML/SciMLTutorials.jl>. For more information on doing scientific machine learning (SciML) with open source software, check out <https://sciml.ai/>.

To locally run this tutorial, do the following commands:

```
using SciMLTutorials
SciMLTutorials.weave_file("introduction", "01-ode_introduction.jmd")
```

Computer Information:

```
Julia Version 1.4.2
Commit 44fa15b150* (2020-05-23 18:35 UTC)
Platform Info:
  OS: Linux (x86_64-pc-linux-gnu)
  CPU: Intel(R) Core(TM) i7-9700K CPU @ 3.60GHz
  WORD_SIZE: 64
  LIBM: libopenlibm
  LLVM: libLLVM-8.0.1 (ORCJIT, skylake)
```

Environment:

```
JULIA_LOAD_PATH = /builds/JuliaGPU/DiffEqTutorials.jl:
JULIA_DEPOT_PATH = /builds/JuliaGPU/DiffEqTutorials.jl/.julia
JULIA_CUDA_MEMORY_LIMIT = 2147483648
JULIA_NUM_THREADS = 8
```

Package Information:

```
Status `~/builds/JuliaGPU/DiffEqTutorials.jl/tutorials/introduction/Project.toml`
[6e4b80f9-dd63-53aa-95a3-0cdb28fa8baf] BenchmarkTools 0.5.0
[0c46a032-eb83-5123-abaf-570d42b7fbaf] DifferentialEquations 6.15.0
[65888b18-ceab-5e60-b2b9-181511a3b968] ParameterizedFunctions 5.6.0
[91a5bcdd-55d7-5caf-9e0b-520d859cae80] Plots 1.6.3
[90137ffa-7385-5640-81b9-e52037218182] StaticArrays 0.12.4
[c3572dad-4567-51f8-b174-8c6c989267f4] Sundials 4.3.0
[37e2e46d-f89d-539d-b4ee-838fcccc9c8e] LinearAlgebra
```