

An Intro to DifferentialEquations.jl

Chris Rackauckas

October 7, 2020

0.1 Basic Introduction Via Ordinary Differential Equations

This notebook will get you started with DifferentialEquations.jl by introducing you to the functionality for solving ordinary differential equations (ODEs). The corresponding documentation page is the [ODE tutorial](#). While some of the syntax may be different for other types of equations, the same general principles hold in each case. Our goal is to give a gentle and thorough introduction that highlights these principles in a way that will help you generalize what you have learned.

0.1.1 Background

If you are new to the study of differential equations, it can be helpful to do a quick background read on [the definition of ordinary differential equations](#). We define an ordinary differential equation as an equation which describes the way that a variable u changes, that is

$$u' = f(u, p, t)$$

where p are the parameters of the model, t is the time variable, and f is the nonlinear model of how u changes. The initial value problem also includes the information about the starting value:

$$u(t_0) = u_0$$

Together, if you know the starting value and you know how the value will change with time, then you know what the value will be at any time point in the future. This is the intuitive definition of a differential equation.

0.1.2 First Model: Exponential Growth

Our first model will be the canonical exponential growth model. This model says that the rate of change is proportional to the current value, and is this:

$$u' = au$$

where we have a starting value $u(0) = u_0$. Let's say we put 1 dollar into Bitcoin which is increasing at a rate of 98% per year. Then calling now $t = 0$ and measuring time in years, our model is:

$$u' = 0.98u$$

and $u(0) = 1.0$. We encode this into Julia by noticing that, in this setup, we match the general form when

```
f(u,p,t) = 0.98u
```

```
f (generic function with 1 method)
```

with `u_0 = 1.0`. If we want to solve this model on a time span from `t=0.0` to `t=1.0`, then we define an `ODEProblem` by specifying this function `f`, this initial condition `u0`, and this time span as follows:

```
using DifferentialEquations
f(u,p,t) = 0.98u
u0 = 1.0
tspan = (0.0,1.0)
prob = ODEProblem(f,u0,tspan)
```

```
ODEProblem with uType Float64 and tType Float64. In-place: false
timespan: (0.0, 1.0)
u0: 1.0
```

To solve our `ODEProblem` we use the command `solve`.


```
sol = solve(prob)
```

```
retcode: Success
Interpolation: automatic order switching interpolation
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

and that's it: we have successfully solved our first ODE!

Analyzing the Solution Of course, the solution type is not interesting in and of itself. We want to understand the solution! The documentation page which explains in detail the functions for analyzing the solution is the [Solution Handling](#) page. Here we will describe some of the basics. You can plot the solution using the plot recipe provided by [Plots.jl](#):


```
using Plots; gr()
plot(sol)
```



From the picture we see that the solution is an exponential curve, which matches our intuition. As a plot recipe, we can annotate the result using any of the [Plots.jl attributes](#). For example:


```
plot(sol,linewidth=5,title="Solution to the linear ODE with a thick line",
      axis="Time (t)",axis="u(t) (in μm)",label="My Thick Line!") # legend=false
```

Solution to the linear ODE with a thick line



Using the mutating `plot!` command we can add other pieces to our plot. For this ODE we know that the true solution is $u(t) = u_0 \exp(at)$, so let's add some of the true solution to our plot:

```
plot!(sol.t, t->1.0*exp(0.98t),lw=3,ls=:dash,label="True Solution!")
```



In the previous command I demonstrated `sol.t`, which grabs the array of time points that the solution was saved at:

```
sol.t

5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
```

We can get the array of solution values using `sol.u`:

```
sol.u

5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

`sol.u[i]` is the value of the solution at time `sol.t[i]`. We can compute arrays of functions of the solution values using standard comprehensions, like:

```
[t+u for (u,t) in tuples(sol)]
```

```
5-element Array{Float64,1}:
 1.0
 1.2038471492789395
 1.7643769243364813
 2.66648203038311
 3.664456142481451
```

However, one interesting feature is that, by default, the solution is a continuous function. If we check the print out again:

```
sol

retcode: Success
Interpolation: automatic order switching interpolation
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

you see that it says that the solution has a order changing interpolation. The default algorithm automatically switches between methods in order to handle all types of problems. For non-stiff equations (like the one we are solving), it is a continuous function of 4th order accuracy. We can call the solution as a function of time `sol(t)`. For example, to get the value at `t=0.45`, we can use the command:

```
sol(0.45)
```

```
1 . 5 5 4 2 6 1 0 4 8 0 5 5 3 1 2
```

Controlling the Solver `DifferentialEquations.jl` has a common set of solver controls among its algorithms which can be found [at the Common Solver Options](#) page. We will detail some of the most widely used options.

The most useful options are the tolerances `abstol` and `reltol` . These tell the internal adaptive time stepping engine how precise of a solution you want. Generally, `reltol` is the relative accuracy while `abstol` is the accuracy when `u` is near zero. These tolerances are local tolerances and thus are not global guarantees. However, a good rule of thumb is that the total solution accuracy is 1-2 digits less than the relative tolerances. Thus for the defaults `abstol=1e-6` and `reltol=1e-3` , you can expect a global accuracy of about 1-2 digits. If we want to get around 6 digits of accuracy, we can use the commands:

```
sol = solve(prob, abstol=1e-8, reltol=1e-8)

retcode: Success
Interpolation: automatic order switching interpolation
t: 9-element Array{Float64,1}:
 0.0
 0.04127492324135852
 0.14679917846877366
```

```

0.28631546412766684
0.4381941361169628
0.6118924302028597
0.7985659100883337
0.9993516479536952
1.0
u: 9-element Array{Float64,1}:
1.0
1.0412786454705882
1.1547261252949712
1.3239095703537043
1.5363819257509728
1.8214895157178692
2.1871396448296223
2.662763824115295
2.664456241933517


```

Now we can see no visible difference against the true solution:

```

plot(sol)
plot!(sol.t, t->1.0*exp(0.98t),lw=3,ls=:dash,label="True Solution!")

```



Notice that by decreasing the tolerance, the number of steps the solver had to take was 9 instead of the previous 5. There is a trade off between accuracy and speed, and it is up to you to determine what is the right balance for your problem.

Another common option is to use `saveat` to make the solver save at specific time points. For example, if we want the solution at an even grid of $t=0.1k$ for integers k , we would use the command:

```

sol = solve(prob,saveat=0.1)

retcode: Success

```

```

Interpolation: 1st order linear
t: 11-element Array{Float64,1}:
 0.0
 0.1
 0.2
 0.3
 0.4
 0.5
 0.6
 0.7
 0.8
 0.9
 1.0
u: 11-element Array{Float64,1}:
 1.0
 1.102962785129292
 1.2165269512238264
 1.341783821227542
 1.4799379510586077
 1.632316207054161
 1.8003833264983584
 1.9857565541588758
 2.1902158127997695
 2.415725742084496
 2.664456142481451

```

Notice that when `saveat` is used the continuous output variables are no longer saved and thus `sol(t)`, the interpolation, is only first order. We can save at an uneven grid of points by passing a collection of values to `saveat`. For example:

```
sol = solve(prob,saveat=[0.2,0.7,0.9])
```

```

retcode: Success
Interpolation: 1st order linear
t: 3-element Array{Float64,1}:
 0.2
 0.7
 0.9
u: 3-element Array{Float64,1}:
 1.2165269512238264
 1.9857565541588758
 2.415725742084496

```

If we need to reduce the amount of saving, we can also turn off the continuous output directly via `dense=false`:

```

sol = solve(prob,dense=false)

retcode: Success
Interpolation: 1st order linear
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448

```

```
1.9730384275622996
2.664456142481451
```

and to turn off all intermediate saving we can use `save_everystep=false`:

```
sol = solve(prob,save_everystep=false)
```

```
retcode: Success
Interpolation: 1st order linear
t: 2-element Array{Float64,1}:
 0.0
 1.0
u: 2-element Array{Float64,1}:
 1.0
 2.664456142481451
```

If we want to solve and only save the final value, we can even set `save_start=false`.

```
sol = solve(prob,save_everystep=false,save_start = false)
```

```
retcode: Success
Interpolation: 1st order linear
t: 1-element Array{Float64,1}:
 1.0
u: 1-element Array{Float64,1}:
 2.664456142481451
```

Note that similarly on the other side there is `save_end=false`.

More advanced saving behaviors, such as saving functionals of the solution, are handled via the `SavingCallback` in the [Callback Library](#) which will be addressed later in the tutorial.

Choosing Solver Algorithms There is no best algorithm for numerically solving a differential equation. When you call `solve(prob)`, `DifferentialEquations.jl` makes a guess at a good algorithm for your problem, given the properties that you ask for (the tolerances, the saving information, etc.). However, in many cases you may want more direct control. A later notebook will help introduce the various *algorithms* in `DifferentialEquations.jl`, but for now let's introduce the *syntax*.

The most crucial determining factor in choosing a numerical method is the stiffness of the model. Stiffness is roughly characterized by a Jacobian `f` with large eigenvalues. That's quite mathematical, and we can think of it more intuitively: if you have big numbers in `f` (like parameters of order `1e5`), then it's probably stiff. Or, as the creator of the MATLAB ODE Suite, Lawrence Shampine, likes to define it, if the standard algorithms are slow, then it's stiff. We will go into more depth about diagnosing stiffness in a later tutorial, but for now note that if you believe your model may be stiff, you can hint this to the algorithm chooser via `alg_hints = [:stiff]`.

```
sol = solve(prob,alg_hints=[:stiff])
```

```
retcode: Success
Interpolation: specialized 3rd order "free" stiffness-aware interpolation
t: 8-element Array{Float64,1}:
 0.0
 0.05653299582822294
 0.17270731152826024
 0.3164602871490142
```



```

0.5057500163821153
0.7292241858994543
0.9912975001018789
1.0
u: 8-element Array{Float64,1}:
1.0
1.0569657840332976
1.1844199383303913
1.3636037723365293
1.6415399686182572
2.0434491434754793
2.641825616057761
2.6644526430553817

```

Stiff algorithms have to solve implicit equations and linear systems at each step so they should only be used when required.

If we want to choose an algorithm directly, you can pass the algorithm type after the problem as `solve(prob,alg)`. For example, let's solve this problem using the `Tsit5()` algorithm, and just for show let's change the relative tolerance to `1e-6` at the same time:

```

sol = solve(prob,Tsit5(),reltol=1e-6)

retcode: Success
Interpolation: specialized 4th order "free" interpolation
t: 10-element Array{Float64,1}:
0.0
0.028970819746309166
0.10049147151547619
0.19458908698515082
0.3071725081673423
0.43945421453622546
0.5883434923759523
0.7524873357619015
0.9293021330536031
1.0
u: 10-element Array{Float64,1}:
1.0
1.0287982807225062
1.1034941463604806
1.2100931078233779
1.351248605624241
1.538280340326815
1.7799346012651116
2.090571742234628
2.486102171447025
2.6644562434913377

```

0.1.3 Systems of ODEs: The Lorenz Equation

Now let's move to a system of ODEs. The [Lorenz equation](#) is the famous "butterfly attractor" that spawned chaos theory. It is defined by the system of ODEs:

$$\frac{dx}{dt} = \sigma(y - x) \quad (1)$$

$$\frac{dy}{dt} = x(\rho - z) - y \quad (2)$$

$$\frac{dz}{dt} = xy - \beta z \quad (3)$$

To define a system of differential equations in `DifferentialEquations.jl`, we define our `f` as a vector function with a vector initial condition. Thus, for the vector `u = [x,y,z]'`, we have the derivative function:

```
function lorenz!(du,u,p,t)
    σ,ρ,β = p
    du[1] = σ*(u[2]-u[1])
    du[2] = u[1]*(ρ-u[3]) - u[2]
    du[3] = u[1]*u[2] - β*u[3]
end
```

```
lorenz! (generic function with 1 method)
```

Notice here we used the in-place format which writes the output to the preallocated vector `du`. For systems of equations the in-place format is faster. We use the initial condition $u_0 = [1.0, 0.0, 0.0]$ as follows:

```
u0 = [1.0,0.0,0.0]
```

```
3-element Array{Float64,1}:
 1.0
 0.0
 0.0
```

Lastly, for this model we made use of the parameters `p`. We need to set this value in the `ODEProblem` as well. For our model we want to solve using the parameters $\sigma = 10$, $\rho = 28$, and $\beta = 8/3$, and thus we build the parameter collection:

```
p = (10,28,8/3) # we could also make this an array, or any other type!

(10, 28, 2.6666666666666665)
```

Now we generate the `ODEProblem` type. In this case, since we have parameters, we add the parameter values to the end of the constructor call. Let's solve this on a time span of `t=0` to `t=100`:

```
tspan = (0.0,100.0)
prob = ODEProblem(lorenz!,u0,tspan,p)
```

```
ODEProblem with uType Array{Float64,1} and tType Float64. In-place: true
timespan: (0.0, 100.0)
u0: [1.0, 0.0, 0.0]
```

Now, just as before, we solve the problem:

```
sol = solve(prob)

retcode: Success
Interpolation: automatic order switching interpolation
t: 1294-element Array{Float64,1}:
```

```

0.0
3.5678604836301404e-5
0.0003924646531993154
0.0032624077544510573
0.009058075635317072
0.01695646895607931
0.0276899566248403
0.041856345938267966
0.06024040228733675
0.08368539694547242

:*(99.3940307091529799.4700114749437599.5437965690901599.61465155834999.6909382314810199.787330232337
1294-element Array{Array{Float64,1},1}:
 [1.0, 0.0, 0.0]
 [0.9996434557625105, 0.0009988049817849058, 1.781434788799208e-8]
 [0.9961045497425811, 0.010965399721242457, 2.146955365838907e-6]
 [0.9693591634199452, 0.08977060667778931, 0.0001438018342266937]
 [0.9242043615038835, 0.24228912482984957, 0.0010461623302512404]
 [0.8800455868998046, 0.43873645009348244, 0.0034242593451028745]
 [0.8483309877783048, 0.69156288756671, 0.008487623500490047]
 [0.8495036595681027, 1.0145425335433382, 0.01821208597613427]
 [0.9139069079152129, 1.4425597546855036, 0.03669381053327124]
 [1.0888636764765296, 2.052326153029042, 0.07402570506414284]

:*( [12.999157033749652, 14.10699925404482, 31.74244844521858] [11.646131422021162,
7.2855792145502845, 35.365000488215486] [7.777555445486692, 2.5166095828739574,
32.030953593541675] [4.739741627223412, 1.5919220588229062,
27.249779003951755] [3.2351668945618774, 2.3121727966182695,
22.724936101772805] [3.310411964698304, 4.28106626744641,
18.435441144016366] [4.527117863517627, 6.895878639772805,
16.58544600757436] [8.043672261487556, 12.711555298531689,
18.12537420595938] [9.97537965430362, 15.143884806010783, 21.00643286956427]

```

The same solution handling features apply to this case. Thus `sol.t` stores the time points and `sol.u` is an array storing the solution at the corresponding time points.

However, there are a few extra features which are good to know when dealing with systems of equations. First of all, `sol` also acts like an array. `sol[i]` returns the solution at the *i*th time point.

```

sol.t[10],sol[10]

(0.08368539694547242, [1.0888636764765296, 2.052326153029042, 0.07402570506
414284])

```

Additionally, the solution acts like a matrix where `sol[j,i]` is the value of the *j*th variable at time *i*:

```

sol[2,10]

2.052326153029042

```

We can get a real matrix by performing a conversion:

```


A = Array(sol)

3×1294 Array{Float64,2}:
 1.0  0.999643  0.996105  0.969359  ...*( 4.52712 8.04367 9.975380.0 0.000998805
0.0109654 0.0897706 6.89588 12.7116 15.14390.0 1.78143e-8 2.14696e-6 0.000143802 16.5854
18.1254 21.0064

```


This is the same as `sol`, i.e. `sol[i,j] = A[i,j]`, but now it's a true matrix. Plotting will by default show the time series for each variable:

```
plot(sol)
```




If we instead want to plot values against each other, we can use the `vars` command. Let's plot variable 1 against variable 2 against variable 3:

```
plot(sol,vars=(1,2,3))
```




This is the classic Lorenz attractor plot, where the x axis is $u[1]$, the y axis is $u[2]$, and the z axis is $u[3]$. Note that the plot recipe by default uses the interpolation, but we can turn this off:

```
plot(sol, vars=(1,2,3), denseplot=false)
```



Yikes! This shows how calculating the continuous solution has saved a lot of computational effort by computing only a sparse solution and filling in the values! Note that in vars, 0=time, and thus we can plot the time series of a single component like:

```
plot(sol, vars=(0,2))
```



0.2 Internal Types

The last basic user-interface feature to explore is the choice of types. DifferentialEquations.jl respects your input types to determine the internal types that are used. Thus since in the previous cases, when we used `Float64` values for the initial condition, this meant that the internal values would be solved using `Float64`. We made sure that time was specified via `Float64` values, meaning that time steps would utilize 64-bit floats as well. But, by simply changing these types we can change what is used internally.

As a quick example, let's say we want to solve an ODE defined by a matrix. To do this, we can simply use a matrix as input.

```
A = [1. 0 0 -5
      4 -2 4 -3
      -4 0 0 1
      5 -2 2 3]
u0 = rand(4,2)
tspan = (0.0,1.0)
f(u,p,t) = A*u
prob = ODEProblem(f,u0,tspan)
sol = solve(prob)
```

```
retcode: Success
Interpolation: automatic order switching interpolation
```

```

t: 10-element Array{Float64,1}:
 0.0
 0.05146877813401059
 0.1282527639612372
 0.20731579115404317
 0.31194657697845707
 0.4225431598147541
 0.5523447290123698
 0.6929130710324516
 0.8390639277681302
 1.0
u: 10-element Array{Array{Float64,2},1}:
 [0.9338218261404168 0.25027683696626934; 0.6880113011089077 0.506902953589
 7485; 0.28992751685776486 0.29603372052070287; 0.2189820604968442 0.2920193
 8310634223]
 [0.8948426636137807 0.17525069693742912; 0.7903064122469458 0.504947202325
 4512; 0.11788938369023497 0.2690130051528156; 0.45134812084245574 0.3755129
 5336193646]
 [0.7193322867237658 0.017848328300158783; 0.7787283709167452 0.44462855940
 05449; -0.08596618899107833 0.27104811527568967; 0.7872952645502843 0.48037
 868353315794]
 [0.38913424909065064 -0.19495121518957492; 0.5763782165401627 0.3271187297
 261934; -0.19022504229055434 0.3390385836921268; 1.1033190209303758 0.55623
 67883714682]
 [-0.272987172852885 -0.5367162625860921; 0.06637250786299154 0.12676175904
 182013; -0.08904908575802231 0.5510289887642261; 1.4406242592723904 0.59227
 6076515206]
 [-1.2152623883063445 -0.9339362563237532; -0.6552725287351174 -0.069600095
 88756283; 0.4044789366199196 0.9391212647236922; 1.6463184935695707 0.53220
 70391817578]
 [-2.5321196966566832 -1.3667942147334118; -1.4967473770474502 -0.156239524
 16703127; 1.5859245470209573 1.5967175318885862; 1.6099302942619835 0.30950
 920955914035]
 [-3.984027621234267 -1.656187624294827; -1.9947474758819008 0.095607008001
 0418; 3.622513877511197 2.474070467842447; 1.128746260031665 -0.13564013502
 98888]
 [-5.123308606501874 -1.5618097576166066; -1.5152890998010435 0.91571474211
 66967; 6.408345924598824 3.3725094942512968; 0.04087731885134804 -0.8180666
 081388397]
 [-5.284369363948497 -0.7291297770767666; 0.8351113144423721 2.598699526403
 812; 9.701259712382853 3.952168390619043; -1.9011949761480538 -1.7635030272
 350947]

```

There is no real difference from what we did before, but now in this case u_0 is a 4×2 matrix. Because of that, the solution at each time point is matrix:

```
sol[3]
```

```

4×2 Array{Float64,2}:
 0.719332  0.0178483
 0.778728  0.444629
-0.0859662 0.271048
 0.787295  0.480379

```

In `DifferentialEquations.jl`, you can use any type that defines $+$, $-$, $*$, $/$, and has an appropriate `norm`. For example, if we want arbitrary precision floating point numbers, we can change the input to be a matrix of `BigFloat`:

```
big_u0 = big.(u0)
```

```

4×@*(2 Array{*@{BigFloat,2}:
0.933822 0.250277
0.688011 0.506903
0.289928 0.296034
0.218982 0.292019

```

and we can solve the ODEProblem with arbitrary precision numbers by using that initial condition:

```

prob = ODEProblem(f,big_u0,tspan)
sol = solve(prob)

```

```
retcode: Success
```

```
Interpolation: automatic order switching interpolation
```

```
t: 6-element Array{Float64,1}:

```

```

0.0
0.11057493353303068
0.3511651471634287
0.6079979861780698
0.9128674287985695
1.0

```

```
u: 6-element Array{Array{BigFloat,2},1}:

```

```

[0.9338218261404167552797161988564766943454742431640625 0.2502768369662693
448418622210738249123096466064453125; 0.68801130110890773572407397296046838
16432952880859375 0.5069029535897484617379404880921356379985809326171875; 0
.2899275168577648553736025860416702926158905029296875 0.2960337205207028699
049942588317207992076873779296875; 0.21898206049684421081735763436881825327
87322998046875 0.2920193831063422340577062641386874020099639892578125]

```

```

[0.77239873855741321661606473648931765961927434321684729537657211300378578
15566487 0.0586802681616850825750892044486655011067660232695755494823714494
9459332515385472; 0.7983190269569502508789002441004863061702450516880161534
719197867133224947106807 0.463935237012570159919804950767369093133943465787
6696341118810725997303496111546; -0.046432791404165707748623095400109635169
83236852924305016851196377651592319638598 0.2654669516080728216362141912434
290841978493229115504968168466254059782595683244; 0.71181836663602067519214
67726109225449577557040253106473278323022438069618878704 0.4586651033315070
868834186488585127140115800036273206608866245665779243815812106]

```

```

[-0.5816638934039496850572749036153737301218416026805956924787419713546080
000067551 -0.67597983425871806741602434543339798757174875847306851367220840
4362071844060415; -0.176149192233094996123312756917445800759995206936687842
9379335298818749046650894 0.05088016651512754861025101628800496084126104181
938104016200949981501608398574898; 0.03599237459824926817205771925798226767
495606097667100055714329209569163091010883 0.669199643771909933370283936856
3770515761561277795095694616043664519911641985837; 1.5342850506967671243842
09557714889005644497001216508001344895474686817573808192 0.5835063469846612
714405850427245194743221771964413515699264024033364016453801904]

```

```

[-3.1211284189176428560244364438691582614081312179971082950826992258684312
94454187 -1.513071111332936026161449734212645012165404245427658564772037987
709652991099251; -1.7719946886210192489022833052872999653928053409521773729
03270787160217910628063 -0.109425676668356804426032007995659395714133588848
4146466416331419257999455290733; 2.3015666072876793085539325970957632717982
25019523975523555697958980163267235734 1.9310879221240217062135028484851682
90951759305442921745234754028362373521067534; 1.480173141110081203763672651
9104205250855945236363327627599305446635125460224 0.15910013609599545690896
21908926812031916393379265199586134693292460299475142873]

```

```

[-5.3863033282230495657882179654254135890873746166996718046373705729081617
46065015 -1.290815387092587891861109872912716870469051788637872937647154002
765630916868849; -0.7102367491366308245201787312045460680222313722451027810
676710014284176645176509 1.584413436954855058756533182555749764445098208042
220750232325894191522422994462; 7.94108353179685041132344642151937192614204

```



```

5110365290231595083421616489906527109 3.72237899556946426432155783942557240
7212804119728732857413506299449999276828452; -0.754205719191458953191834502
058314459822854052802833371987597297191908624485419 -1.2334865328087742843
92798712214270301518285989172823010330116176101772820750069]
[-5.2843718170745453627594806414479785710755588685166059644495300817256896
84364062 -0.729125757399840225362707661322376109637650754250178335206557470
2509714421270548; 0.8351258792055375755805729897231989563720499242580608462
50342247722689565971471 2.5987097772570667742186116099342988176198372646893
00435925101049272157344734321; 9.701279433034310021289888233497294429262382
902009906421067438237990814612732321 3.952172372578267379676806534923306786
534931155805109468337046923668578119284993; -1.9012047067893348097771837415
6054003856396075969144047448196919182601235106419 -1.7635082000820369760165
22737979390998358223539291664992244896286197180448650274]

```

```
sol[1,3]
```

```

- 0 . 5 8 1 6 6 3 8 9 3 4 0 3 9 4 9 6 8 5 0 5 7 2 7 4 9 0 3 6 1 5 3 7 3 7 3 0 1 2 1 8 4 1 6 0 2 6
8 0 5 9 5 6 9 2 4 7 8 7 4 1 9 7 1 3 5 4 6 0 8 0 0 0 0 0 6 7 5 5 1

```

To really make use of this, we would want to change `abstol` and `reltol` to be small! Notice that the type for "time" is different than the type for the dependent variables, and this can be used to optimize the algorithm via keeping multiple precisions. We can convert time to be arbitrary precision as well by defining our time span with `BigFloat` variables:

```

prob = ODEProblem(f,big_u0,big.(tspan))
sol = solve(prob)

```

```
retcode: Success
```

```
Interpolation: automatic order switching interpolation
```

```
t: 6-element Array{BigFloat,1}:
```

```

0.0
0.110574933533030676767044992291410460736647132073819704392974371361429114
4588063
0.351165147163428749606749314916586915591934717778043386039962049618655145
726078
0.607997986178069798527417081826370059196369529289515457261068861808800327
8252208
0.912867428798569551772237310099552470220402614195396826496939216985108284
154487
1.0

```

```
u: 6-element Array{Array{BigFloat,2},1}:
```

```

[0.9338218261404167552797161988564766943454742431640625 0.2502768369662693
448418622210738249123096466064453125; 0.68801130110890773572407397296046838
16432952880859375 0.5069029535897484617379404880921356379985809326171875; 0
.2899275168577648553736025860416702926158905029296875 0.2960337205207028699
049942588317207992076873779296875; 0.21898206049684421081735763436881825327
87322998046875 0.2920193831063422340577062641386874020099639892578125]
[0.77239873855741321281565563711434221209855972643192570919415972792619720
09619641 0.0586802681616850795275465264844029171261532413860976973210435732
509603409422369; 0.79831902695695024974938568727453929695918288015373546460
54843421323922317052816 0.4639352370125701585461106883987053968119638701149
677746995522723882007339667809; -0.0464327914041657109913640334547738146569
5812949594731957662208989264851688732784 0.26546695160807282194162239482887
70624011162947519684356033013409257841160628962; 0.711818366636020681067194
1547489126036106026739371138217110541581467796039222019 0.45866510333150708
86187618324812081640279844200163632378379402391388166324210852]
[-0.5816638934039497668505595000177779345710555592708976780584649659339069
083517362 -0.67597983425871810302997550374666699558451185002393730127418798
70880738691857099; -0.17614919223309505988077975770844472746966157135604333

```

```

71926156536866243125355542 0.0508801665151275299842183353567870576710236109
8343626881554175746017851744897418; 0.0359923745982493064363987537995935957
7271015843771704906076478928926754801291528 0.66919964377190996595073230060
71587047239162596209589878791025008762566306674687; 1.534285050696767145383
03153166407380764745790104508310855308516071321400513083 0.5835063469846612
675482738408630815091731767945588199229826259680228167366278924]
[-3.1211284189176429953062270034967107741998995499052761154609154315443887
34010964 -1.513071111332936056720486902199187243591471615925701885344744119
423369147077328; -1.7719946886210193041646640899352064701813630180714796425
49590772958119657237336 -0.109425676668356785713678046705182763321444210359
605457083164059287946420955607; 2.30156660728767949340735207976614250943583
8544111146437360865856976463084821872 1.93108792212402178843487233509929583
5141442588864691067256741929613632530326687; 1.4801731411100811638137870472
01537883711889469583177238109640148405776516967185 0.1591001360959954171043
172107343855158018475235586113665691748337176267667579001]
[-5.3863033282230495872703962564851539035219766567197892121439981197296255
59922762 -1.290815387092587827005039547768160000273194195159802363442007475
83568641195931; -0.71023674913663063962914661879953702391877654185491772150
68313803970601057077066 1.5844134369548551951802478805986849086229411316205
63955265606571309420744232915; 7.941083531796850687831310488201025632631650
268148628234558797237950784509390913 3.722378995569464316585195039331591425
7485848774724566259961001497935145419303; -0.754205719191459111341525292284
6254680592856194351080839124157678067861345567135 -1.2334865328087743625747
06340038551977873895553056894263817666111317576780329665]
[-5.2843718170745454213466655274735834620264261490050532513161592337135297
8735029 -0.7291257573998403376120177645344218597108177905824524404845661636
910427367813847; 0.83512587920553727441879952382991487581037075155285346678
3690648364765834934648 2.59870977725706659401149073835806331310301987934332
5907536412354476659842138159; 9.7012794330343097543328373464788896721711821
15897980798085563498657608294984934 3.9521723725782673636757874706321707646
59475258281353348670424409216177736269991; -1.90120470678933461003187819868
6296574953366886646848575174567781235554753161944 -1.7635082000820368895686
15084575900657566627477278290997446678749421502038705135]

```

Let's end by showing a more complicated use of types. For small arrays, it's usually faster to do operations on static arrays via the package [StaticArrays.jl](#). The syntax is similar to that of normal arrays, but for these special arrays we utilize the `@SMatrix` macro to indicate we want to create a static array.

```

using StaticArrays
A = @SMatrix [ 1.0  0.0 0.0 -5.0
               4.0 -2.0 4.0 -3.0
              -4.0  0.0 0.0  1.0
               5.0 -2.0 2.0  3.0]
u0 = @SMatrix rand(4,2)
tspan = (0.0,1.0)
f(u,p,t) = A*u
prob = ODEProblem(f,u0,tspan)
sol = solve(prob)

retcode: Success
Interpolation: automatic order switching interpolation
t: 12-element Array{Float64,1}:
 0.0
 0.01128478320136641
 0.04432586691283831
 0.1006982762581945
 0.1707099286973292
 0.255116869837141

```

```

0.35338226190049626
0.4792279457595038
0.6232734912548586
0.7769167439253583
0.9425102420651629
1.0
u: 12-element Array{StaticArrays.SArray{Tuple{4,2},Float64,2,8},1}:
 [0.3207548495240642 0.8958480860767932; 0.05145889888617261 0.463421267499
63297; 0.7617439410311491 0.011218901523266211; 0.7149661509423113 0.607770
8489159379]
 [0.2821808873524267 0.8698055278320413; 0.07274078864557351 0.470889734075
241; 0.7565178980737441 -0.02144342377855661; 0.7729001066888023 0.66853448
89058307]
 [0.14799189452720288 0.7720115937701326; 0.11013857521294577 0.46500569436
952344; 0.7560196542049209 -0.10529216495518895; 0.9373025225069743 0.84416
12987564938]
 [-0.15298478118424083 0.530261321928576; 0.0937020971578775 0.362997545605
63124; 0.8151500049895408 -0.19811650473157774; 1.1948583008078497 1.131754
4601518477]
 [-0.6470185470615244 0.09894436058777056; -0.04708620891365348 0.089581214
20544468; 1.017543669945439 -0.19873737519735588; 1.4604617143392435 1.4556
789185454904]
 [-1.3981984202990074 -0.6056311495767073; -0.3403349806960074 -0.414028076
9664418; 1.491675288341294 0.01826623939914243; 1.67313935255188 1.77246727
10082308]
 [-2.430195268099975 -1.6479738035463514; -0.7363369574880048 -1.1430303311
050458; 2.4088462895729705 0.6408571670291584; 1.727491249897271 1.99683194
68649864]
 [-3.8375532555960423 -3.2222844779981794; -1.0560404439682476 -2.056525446
782232; 4.189761002242042 2.1128749351847858; 1.4121971533517272 1.97263538
131972]
 [-5.195026884126617 -5.069506694347904; -0.6504154509301505 -2.58058075975
38326; 6.957192800911287 4.75883152208732; 0.4115677673020941 1.38106006313
05563]
 [-5.692173737110967 -6.536562249871675; 1.3878905551898326 -1.780397632094
5232; 10.30187225351298 8.483102536572735; -1.4835205569181809 -0.061474574
136404]
 [-4.162349577465936 -6.611270586777157; 6.119459024595115 1.54557509062384
26; 13.236546073136779 12.74550520847462; -4.402701390200441 -2.63023565159
13834]
 [-2.9364381041711614 -6.062757991968209; 8.433589842386493 3.4299216604888
72; 13.774156521541126 14.026629420549309; -5.569041351252945 -3.7502670408
7143]

sol[3]

4×@*(2 StaticArrays.SArray{Tuple{4,2},Float64,2,8} with indices SOneTo(4)
×@*(SOneTo(2):0.147992 0.7720120.110139 0.4650060.75602 -0.1052920.937303
0.844161

```

0.3 Conclusion

These are the basic controls in `DifferentialEquations.jl`. All equations are defined via a problem type, and the `solve` command is used with an algorithm choice (or the default) to get a solution. Every solution acts the same, like an array `sol[i]` with `sol.t[i]`, and also like a continuous function `sol(t)` with a nice plot command `plot(sol)`. The Common Solver Options can be used to control the solver for any equation type. Lastly, the types

used in the numerical solving are determined by the input types, and this can be used to solve with arbitrary precision and add additional optimizations (this can be used to solve via GPUs for example!). While this was shown on ODEs, these techniques generalize to other types of equations as well.

0.4 Appendix

This tutorial is part of the SciMLTutorials.jl repository, found at: <https://github.com/SciML/SciMLTutorials>. For more information on doing scientific machine learning (SciML) with open source software, check out <https://sciml.ai/>.

To locally run this tutorial, do the following commands:

```
using SciMLTutorials
SciMLTutorials.weave_file("introduction", "01-ode_introduction.jmd")
```

Computer Information:

```
Julia Version 1.4.2
Commit 44fa15b150* (2020-05-23 18:35 UTC)
Platform Info:
  OS: Linux (x86_64-pc-linux-gnu)
  CPU: Intel(R) Core(TM) i7-9700K CPU @ 3.60GHz
  WORD_SIZE: 64
  LIBM: libopenlibm
  LLVM: libLLVM-8.0.1 (ORCJIT, skylake)
Environment:
  JULIA_LOAD_PATH = /builds/JuliaGPU/DiffEqTutorials.jl:
  JULIA_DEPOT_PATH = /builds/JuliaGPU/DiffEqTutorials.jl/.julia
  JULIA_CUDA_MEMORY_LIMIT = 2147483648
  JULIA_NUM_THREADS = 8
```

Package Information:

```
Status `~/builds/JuliaGPU/DiffEqTutorials.jl/tutorials/introduction/Project.toml`
[6e4b80f9-dd63-53aa-95a3-0cdb28fa8baf] BenchmarkTools 0.5.0
[0c46a032-eb83-5123-abaf-570d42b7fbaf] DifferentialEquations 6.15.0
[65888b18-ceab-5e60-b2b9-181511a3b968] ParameterizedFunctions 5.6.0
[91a5bcdd-55d7-5caf-9e0b-520d859cae80] Plots 1.6.9
[90137ffa-7385-5640-81b9-e52037218182] StaticArrays 0.12.4
[c3572dad-4567-51f8-b174-8c6c989267f4] Sundials 4.3.0
[37e2e46d-f89d-539d-b4ee-838fcccc9c8e] LinearAlgebra
```