

An Intro to DifferentialEquations.jl

Chris Rackauckas

July 28, 2020

0.1 Basic Introduction Via Ordinary Differential Equations

This notebook will get you started with DifferentialEquations.jl by introducing you to the functionality for solving ordinary differential equations (ODEs). The corresponding documentation page is the [ODE tutorial](#). While some of the syntax may be different for other types of equations, the same general principles hold in each case. Our goal is to give a gentle and thorough introduction that highlights these principles in a way that will help you generalize what you have learned.

0.1.1 Background

If you are new to the study of differential equations, it can be helpful to do a quick background read on [the definition of ordinary differential equations](#). We define an ordinary differential equation as an equation which describes the way that a variable u changes, that is

$$u' = f(u, p, t)$$

where p are the parameters of the model, t is the time variable, and f is the nonlinear model of how u changes. The initial value problem also includes the information about the starting value:

$$u(t_0) = u_0$$

Together, if you know the starting value and you know how the value will change with time, then you know what the value will be at any time point in the future. This is the intuitive definition of a differential equation.

0.1.2 First Model: Exponential Growth

Our first model will be the canonical exponential growth model. This model says that the rate of change is proportional to the current value, and is this:

$$u' = au$$

where we have a starting value $u(0) = u_0$. Let's say we put 1 dollar into Bitcoin which is increasing at a rate of 98% per year. Then calling now $t = 0$ and measuring time in years, our model is:

$$u' = 0.98u$$

and $u(0) = 1.0$. We encode this into Julia by noticing that, in this setup, we match the general form when

```
f(u,p,t) = 0.98u
```

```
f (generic function with 1 method)
```

with `u_0 = 1.0`. If we want to solve this model on a time span from `t=0.0` to `t=1.0`, then we define an `ODEProblem` by specifying this function `f`, this initial condition `u0`, and this time span as follows:

```
using DifferentialEquations
f(u,p,t) = 0.98u
u0 = 1.0
tspan = (0.0,1.0)
prob = ODEProblem(f,u0,tspan)
```

```
ODEProblem with uType Float64 and tType Float64. In-place: false
timespan: (0.0, 1.0)
u0: 1.0
```

To solve our `ODEProblem` we use the command `solve`.


```
sol = solve(prob)
```

```
retcode: Success
Interpolation: automatic order switching interpolation
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

and that's it: we have successfully solved our first ODE!

Analyzing the Solution Of course, the solution type is not interesting in and of itself. We want to understand the solution! The documentation page which explains in detail the functions for analyzing the solution is the [Solution Handling](#) page. Here we will describe some of the basics. You can plot the solution using the plot recipe provided by [Plots.jl](#):


```
using Plots; gr()
plot(sol)
```



From the picture we see that the solution is an exponential curve, which matches our intuition. As a plot recipe, we can annotate the result using any of the [Plots.jl attributes](#). For example:

```
plot(sol,linewidth=5,title="Solution to the linear ODE with a thick line",
      axis="Time (t)",axis="u(t) (in  $\mu\text{m}$ )",label="My Thick Line!") # legend=false
```


Solution to the linear ODE with a thick line



Using the mutating `plot!` command we can add other pieces to our plot. For this ODE we know that the true solution is $u(t) = u_0 \exp(at)$, so let's add some of the true solution to our plot:

```
plot!(sol.t, t->1.0*exp(0.98t),lw=3,ls=:dash,label="True Solution!")
```

Solution to the linear ODE with a thick line



In the previous command I demonstrated `sol.t`, which grabs the array of time points that the solution was saved at:

```
sol.t
```

```
5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
```

We can get the array of solution values using `sol.u`:

```
sol.u
```

```
5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

`sol.u[i]` is the value of the solution at time `sol.t[i]`. We can compute arrays of functions of the solution values using standard comprehensions, like:

```
[t+u for (u,t) in tuples(sol)]
```

```
5-element Array{Float64,1}:
 1.0
 1.2038471492789395
 1.7643769243364813
 2.66648203038311
 3.664456142481451
```

However, one interesting feature is that, by default, the solution is a continuous function. If we check the print out again:

```
sol

retcode: Success
Interpolation: automatic order switching interpolation
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448
 1.9730384275622996
 2.664456142481451
```

you see that it says that the solution has a order changing interpolation. The default algorithm automatically switches between methods in order to handle all types of problems. For non-stiff equations (like the one we are solving), it is a continuous function of 4th order accuracy. We can call the solution as a function of time `sol(t)`. For example, to get the value at `t=0.45`, we can use the command:

```
sol(0.45)

1.554261048055312
```

Controlling the Solver `DifferentialEquations.jl` has a common set of solver controls among its algorithms which can be found [at the Common Solver Options](#) page. We will detail some of the most widely used options.

The most useful options are the tolerances `abstol` and `reltol` . These tell the internal adaptive time stepping engine how precise of a solution you want. Generally, `reltol` is the relative accuracy while `abstol` is the accuracy when `u` is near zero. These tolerances are local tolerances and thus are not global guarantees. However, a good rule of thumb is that the total solution accuracy is 1-2 digits less than the relative tolerances. Thus for the defaults `abstol=1e-6` and `reltol=1e-3` , you can expect a global accuracy of about 1-2 digits. If we want to get around 6 digits of accuracy, we can use the commands:

```
sol = solve(prob, abstol=1e-8, reltol=1e-8)

retcode: Success
Interpolation: automatic order switching interpolation
t: 9-element Array{Float64,1}:
 0.0
 0.04127492324135852
 0.14679917846877366
```

```

0.28631546412766684
0.4381941361169628
0.6118924302028597
0.7985659100883337
0.9993516479536952
1.0
u: 9-element Array{Float64,1}:
1.0
1.0412786454705882
1.1547261252949712
1.3239095703537043
1.5363819257509728
1.8214895157178692
2.1871396448296223
2.662763824115295
2.664456241933517


```

Now we can see no visible difference against the true solution:

```

plot(sol)
plot!(sol.t, t->1.0*exp(0.98t),lw=3,ls=:dash,label="True Solution!")

```



Notice that by decreasing the tolerance, the number of steps the solver had to take was 9 instead of the previous 5. There is a trade off between accuracy and speed, and it is up to you to determine what is the right balance for your problem.

Another common option is to use `saveat` to make the solver save at specific time points. For example, if we want the solution at an even grid of $t=0.1k$ for integers k , we would use the command:

```

sol = solve(prob,saveat=0.1)

retcode: Success

```

```

Interpolation: 1st order linear
t: 11-element Array{Float64,1}:
 0.0
 0.1
 0.2
 0.3
 0.4
 0.5
 0.6
 0.7
 0.8
 0.9
 1.0
u: 11-element Array{Float64,1}:
 1.0
 1.102962785129292
 1.2165269512238264
 1.341783821227542
 1.4799379510586077
 1.632316207054161
 1.8003833264983584
 1.9857565541588758
 2.1902158127997695
 2.415725742084496
 2.664456142481451

```

Notice that when `saveat` is used the continuous output variables are no longer saved and thus `sol(t)`, the interpolation, is only first order. We can save at an uneven grid of points by passing a collection of values to `saveat`. For example:

```
sol = solve(prob,saveat=[0.2,0.7,0.9])
```

```

retcode: Success
Interpolation: 1st order linear
t: 3-element Array{Float64,1}:
 0.2
 0.7
 0.9
u: 3-element Array{Float64,1}:
 1.2165269512238264
 1.9857565541588758
 2.415725742084496

```

If we need to reduce the amount of saving, we can also turn off the continuous output directly via `dense=false`:

```

sol = solve(prob,dense=false)

retcode: Success
Interpolation: 1st order linear
t: 5-element Array{Float64,1}:
 0.0
 0.10042494449239292
 0.35218603951893646
 0.6934436028208104
 1.0
u: 5-element Array{Float64,1}:
 1.0
 1.1034222047865465
 1.4121908848175448

```

```
1.9730384275622996
2.664456142481451
```

and to turn off all intermediate saving we can use `save_everystep=false`:

```
sol = solve(prob,save_everystep=false)
```

```
retcode: Success
Interpolation: 1st order linear
t: 2-element Array{Float64,1}:
 0.0
 1.0
u: 2-element Array{Float64,1}:
 1.0
 2.664456142481451
```

If we want to solve and only save the final value, we can even set `save_start=false`.

```
sol = solve(prob,save_everystep=false,save_start = false)
```

```
retcode: Success
Interpolation: 1st order linear
t: 1-element Array{Float64,1}:
 1.0
u: 1-element Array{Float64,1}:
 2.664456142481451
```

Note that similarly on the other side there is `save_end=false`.

More advanced saving behaviors, such as saving functionals of the solution, are handled via the `SavingCallback` in the [Callback Library](#) which will be addressed later in the tutorial.

Choosing Solver Algorithms There is no best algorithm for numerically solving a differential equation. When you call `solve(prob)`, `DifferentialEquations.jl` makes a guess at a good algorithm for your problem, given the properties that you ask for (the tolerances, the saving information, etc.). However, in many cases you may want more direct control. A later notebook will help introduce the various *algorithms* in `DifferentialEquations.jl`, but for now let's introduce the *syntax*.

The most crucial determining factor in choosing a numerical method is the stiffness of the model. Stiffness is roughly characterized by a Jacobian \mathbf{f} with large eigenvalues. That's quite mathematical, and we can think of it more intuitively: if you have big numbers in \mathbf{f} (like parameters of order $1e5$), then it's probably stiff. Or, as the creator of the MATLAB ODE Suite, Lawrence Shampine, likes to define it, if the standard algorithms are slow, then it's stiff. We will go into more depth about diagnosing stiffness in a later tutorial, but for now note that if you believe your model may be stiff, you can hint this to the algorithm chooser via `alg_hints = [:stiff]`.

```
sol = solve(prob,alg_hints=[:stiff])
```

```
retcode: Success
Interpolation: specialized 3rd order "free" stiffness-aware interpolation
t: 8-element Array{Float64,1}:
 0.0
 0.05653299582822294
 0.17270731152826024
 0.3164602871490142
```



```

0.5057500163821153
0.7292241858994543
0.9912975001018789
1.0
u: 8-element Array{Float64,1}:
1.0
1.0569657840332976
1.1844199383303913
1.3636037723365293
1.6415399686182572
2.0434491434754793
2.641825616057761
2.6644526430553817

```

Stiff algorithms have to solve implicit equations and linear systems at each step so they should only be used when required.

If we want to choose an algorithm directly, you can pass the algorithm type after the problem as `solve(prob,alg)`. For example, let's solve this problem using the `Tsit5()` algorithm, and just for show let's change the relative tolerance to `1e-6` at the same time:

```

sol = solve(prob,Tsit5(),reltol=1e-6)

retcode: Success
Interpolation: specialized 4th order "free" interpolation
t: 10-element Array{Float64,1}:
0.0
0.028970819746309166
0.10049147151547619
0.19458908698515082
0.3071725081673423
0.43945421453622546
0.5883434923759523
0.7524873357619015
0.9293021330536031
1.0
u: 10-element Array{Float64,1}:
1.0
1.0287982807225062
1.1034941463604806
1.2100931078233779
1.351248605624241
1.538280340326815
1.7799346012651116
2.090571742234628
2.486102171447025
2.6644562434913377

```

0.1.3 Systems of ODEs: The Lorenz Equation

Now let's move to a system of ODEs. The [Lorenz equation](#) is the famous "butterfly attractor" that spawned chaos theory. It is defined by the system of ODEs:

$$\frac{dx}{dt} = \sigma(y - x) \quad (1)$$

$$\frac{dy}{dt} = x(\rho - z) - y \quad (2)$$

$$\frac{dz}{dt} = xy - \beta z \quad (3)$$

To define a system of differential equations in `DifferentialEquations.jl`, we define our `f` as a vector function with a vector initial condition. Thus, for the vector `u = [x,y,z]'`, we have the derivative function:

```
function lorenz!(du,u,p,t)
    σ,ρ,β = p
    du[1] = σ*(u[2]-u[1])
    du[2] = u[1]*(ρ-u[3]) - u[2]
    du[3] = u[1]*u[2] - β*u[3]
end
```

```
lorenz! (generic function with 1 method)
```

Notice here we used the in-place format which writes the output to the preallocated vector `du`. For systems of equations the in-place format is faster. We use the initial condition $u_0 = [1.0, 0.0, 0.0]$ as follows:

```
u0 = [1.0,0.0,0.0]
```

```
3-element Array{Float64,1}:
 1.0
 0.0
 0.0
```

Lastly, for this model we made use of the parameters `p`. We need to set this value in the `ODEProblem` as well. For our model we want to solve using the parameters $\sigma = 10$, $\rho = 28$, and $\beta = 8/3$, and thus we build the parameter collection:

```
p = (10,28,8/3) # we could also make this an array, or any other type!

(10, 28, 2.6666666666666665)
```

Now we generate the `ODEProblem` type. In this case, since we have parameters, we add the parameter values to the end of the constructor call. Let's solve this on a time span of `t=0` to `t=100`:

```
tspan = (0.0,100.0)
prob = ODEProblem(lorenz!,u0,tspan,p)
```

```
ODEProblem with uType Array{Float64,1} and tType Float64. In-place: true
timespan: (0.0, 100.0)
u0: [1.0, 0.0, 0.0]
```

Now, just as before, we solve the problem:

```
sol = solve(prob)

retcode: Success
Interpolation: automatic order switching interpolation
t: 1294-element Array{Float64,1}:
```

```

0.0
3.5678604836301404e-5
0.0003924646531993154
0.0032624077544510573
0.009058075635317072
0.01695646895607931
0.0276899566248403
0.041856345938267966
0.06024040228733675
0.08368539694547242
:
99.39403070915297
99.47001147494375
99.54379656909015
99.614651558349
99.69093823148101
99.78733023233721
99.86114450046736
99.96115759510786
100.0
u: 1294-element Array{Array{Float64,1},1}:
[1.0, 0.0, 0.0]
[0.9996434557625105, 0.0009988049817849058, 1.781434788799208e-8]
[0.9961045497425811, 0.010965399721242457, 2.146955365838907e-6]
[0.9693591634199452, 0.08977060667778931, 0.0001438018342266937]
[0.9242043615038835, 0.24228912482984957, 0.0010461623302512404]
[0.8800455868998046, 0.43873645009348244, 0.0034242593451028745]
[0.8483309877783048, 0.69156288756671, 0.008487623500490047]
[0.8495036595681027, 1.0145425335433382, 0.01821208597613427]
[0.9139069079152129, 1.4425597546855036, 0.03669381053327124]
[1.0888636764765296, 2.052326153029042, 0.07402570506414284]
:
[12.999157033749652, 14.10699925404482, 31.74244844521858]
[11.646131422021162, 7.2855792145502845, 35.365000488215486]
[7.777555445486692, 2.5166095828739574, 32.030953593541675]
[4.739741627223412, 1.5919220588229062, 27.249779003951755]
[3.2351668945618774, 2.3121727966182695, 22.724936101772805]
[3.310411964698304, 4.28106626744641, 18.435441144016366]
[4.527117863517627, 6.895878639772805, 16.58544600757436]
[8.043672261487556, 12.711555298531689, 18.12537420595938]
[9.97537965430362, 15.143884806010783, 21.00643286956427]

```

The same solution handling features apply to this case. Thus `sol.t` stores the time points and `sol.u` is an array storing the solution at the corresponding time points.

However, there are a few extra features which are good to know when dealing with systems of equations. First of all, `sol` also acts like an array. `sol[i]` returns the solution at the *i*th time point.

```

sol.t[10],sol[10]

(0.08368539694547242, [1.0888636764765296, 2.052326153029042, 0.07402570506
414284])

```

Additionally, the solution acts like a matrix where `sol[j,i]` is the value of the *j*th variable at time *i*:

```

sol[2,10]

```

2.052326153029042

We can get a real matrix by performing a conversion:


```
A = Array(sol)
```

```
3×1294 Array{Float64,2}:
```

```
1.0  0.999643  0.996105  0.969359  ...  4.52712  8.04367  9.97538
0.0  0.000998805 0.0109654 0.0897706  ...  6.89588 12.7116 15.1439
0.0  1.78143e-8 2.14696e-6 0.000143802  ... 16.5854 18.1254 21.0064
```


This is the same as `sol`, i.e. `sol[i,j] = A[i,j]`, but now it's a true matrix. Plotting will by default show the time series for each variable:

```
plot(sol)
```




If we instead want to plot values against each other, we can use the `vars` command. Let's plot variable 1 against variable 2 against variable 3:

```
plot(sol,vars=(1,2,3))
```




This is the classic Lorenz attractor plot, where the x axis is $u[1]$, the y axis is $u[2]$, and the z axis is $u[3]$. Note that the plot recipe by default uses the interpolation, but we can turn this off:

```
plot(sol, vars=(1,2,3), denseplot=false)
```



Yikes! This shows how calculating the continuous solution has saved a lot of computational effort by computing only a sparse solution and filling in the values! Note that in vars, 0=time, and thus we can plot the time series of a single component like:

```
plot(sol, vars=(0,2))
```



0.2 Internal Types

The last basic user-interface feature to explore is the choice of types. DifferentialEquations.jl respects your input types to determine the internal types that are used. Thus since in the previous cases, when we used `Float64` values for the initial condition, this meant that the internal values would be solved using `Float64`. We made sure that time was specified via `Float64` values, meaning that time steps would utilize 64-bit floats as well. But, by simply changing these types we can change what is used internally.

As a quick example, let's say we want to solve an ODE defined by a matrix. To do this, we can simply use a matrix as input.

```
A = [1. 0 0 -5
      4 -2 4 -3
      -4 0 0 1
      5 -2 2 3]
u0 = rand(4,2)
tspan = (0.0,1.0)
f(u,p,t) = A*u
prob = ODEProblem(f,u0,tspan)
sol = solve(prob)
```

```
retcode: Success
```

```
Interpolation: automatic order switching interpolation
```

```

t: 10-element Array{Float64,1}:
 0.0
 0.0431583364054055
 0.1153143854351763
 0.19765891776727906
 0.30243489494902526
 0.42252610292262877
 0.5574409597445212
 0.7002286201489554
 0.8534607321245409
 1.0
u: 10-element Array{Array{Float64,2},1}:
 [0.3984819530794539 0.9637556818431008; 0.9147620825925182 0.3617522521140
 6525; 0.6267976723471336 0.8271326495338405; 0.8092339752422932 0.955089117
 4747071]
 [0.22001705043861136 0.7554289557732575; 0.8819793562014258 0.461039941018
 5265; 0.6112020571939 0.7266240213731241; 0.9671432081880452 1.319994570880
 869]
 [-0.16809430936651382 0.2115347979002603; 0.7235703793558566 0.39397555237
 553544; 0.6792390606369912 0.6972745062519699; 1.188444298085305 1.88520855
 26201756]
 [-0.7318934827269516 -0.6982058934855276; 0.4255744626283403 0.01107806018
 4708305; 0.9297752436191355 0.9476196735462834; 1.3574911015821431 2.423677
 918513722]
 [-1.585613689553726 -2.248520937952204; -0.029183271841557168 -0.796385494
 9827429; 1.5580103553669167 1.8311908363596512; 1.408027829478332 2.8655823
 25376089]
 [-2.6340523430881877 -4.406959946722383; -0.4470237581814744 -1.8201229602
 413465; 2.730774354780132 3.7705886009900373; 1.1852770914783142 2.90727113
 03547463]
 [-3.6615802691423376 -6.941053615430349; -0.4371406305842422 -2.4506613625
 61908; 4.563499899394147 7.192553793559531; 0.5125173036142324 2.1795193612
 067223]
 [-4.186840426001786 -9.062005469554194; 0.5694137887890667 -1.586245349492
 7955; 6.83444470063553 12.004139127139837; -0.7290037820542522 0.3307350388
 000838]
 [-3.5460304772885167 -9.592000913584338; 3.232909106762804 2.2217307543097
 58; 9.033119728808751 17.660432918400907; -2.6303279546178624 -2.9885088035
 790863]
 [-1.2003357638717973 -7.141462398459231; 7.498435076722952 9.4767166751888
 1; 9.976173736001996 21.993644097140702; -4.824170381902354 -7.327832194928
 521]

```

There is no real difference from what we did before, but now in this case `u0` is a **4x2** matrix. Because of that, the solution at each time point is matrix:

```

sol[3]

4×2 Array{Float64,2}:
-0.168094  0.211535
 0.72357   0.393976
 0.679239  0.697275
 1.18844   1.88521

```

In `DifferentialEquations.jl`, you can use any type that defines `+`, `-`, `*`, `/`, and has an appropriate `norm`. For example, if we want arbitrary precision floating point numbers, we can change the input to be a matrix of `BigFloat`:

```
big_u0 = big.(u0)
```

```

4×2 Array{BigFloat,2}:
 0.398482  0.963756
 0.914762  0.361752
 0.626798  0.827133
 0.809234  0.955089

```

and we can solve the ODEProblem with arbitrary precision numbers by using that initial condition:

```

prob = ODEProblem(f,big_u0,tspan)
sol = solve(prob)

```

```
retcode: Success
```

```
Interpolation: automatic order switching interpolation
```

```
t: 5-element Array{Float64,1}:
```

```

0.0
0.15945699885551542
0.42119133406156783
0.7216495970126032
1.0

```

```
u: 5-element Array{Array{BigFloat,2},1}:
```

```

[0.398481953079453887767158448696136474609375 0.96375568184310078656551468
15697662532329559326171875; 0.914762082592518233070677524665370583534240722
65625 0.361752252114065253607577687944285571575164794921875; 0.626797672347
1335720461183882434852421283721923828125 0.82713264953384046940243479184573
51624965667724609375; 0.809233975242293235297097453440073877573013305664062
5 0.9550891174747071499240291814203374087810516357421875]

```

```

[-0.4559188521276426922557296616877006460775274436467549359441118858654108
960309203 -0.23925412414778976888604263078459408735189049808722876519896737
92749897577090273; 0.575680708005470950360061538948728997815827639747374497
3188203393874350146337665 0.22436991924335785137793240243117366222573359599
82229576395335922475501361542198; 0.788644646767375985755451503976023966707
9294624538441025558823347341651145866438 0.78853508225456409849798808480603
21774699109379376158852697107099043622962496344; 1.291542317333483656961441
90106117802624908105477398956574425429132119622288961 2.1910086758340621732
51673284734873053193549689196903413042620626411370691829232]

```

```

[-2.6226197251523085009384094784655856188291900840671220690982652641602263
77383755 -4.381683442002130841435042976632062032158117017380190393487668654
70991763579304; -0.44396293808566942978278119303153645546431749619491129674
15261542302583438074556 -1.809917286516711861185566359402991537394038210650
31563701584969223287248326382; 2.715157351636516006672059824513575494171073
396449181269145099166522153186981791 3.743246217976098517187392262909977867
52318400103534601829448181480835594381392; 1.189604889849235792318452325298
780127452396297056307420039814007449868236232296 2.910076438109654672919005
687643924220091062262782307098662850535858976051392051]

```

```

[-4.1861253001774563833793950521203641663989428491439209482988818840166691
28136406 -9.273734292428036121118625855795516019594788785103234598021511521
845441909424841; 0.83506057611031913375931628649143840598748170631048698679
06482720847223415604546 -1.257779327864159179456344136655481254228177811278
773630835662244490042179079219; 7.17525503088152160742854614009519606729644
4108126769868034174147207319300346556 12.7930048936307755096933747112660014
3164408544803858234982852037306875644757537; -0.961733886232084340345159478
1944616437341812268803610298955081347588574289923561 -0.0500723641690716471
2023838327389803147280102438302008672476191646188377723705673]

```

```

[-1.2003236055657163402902091266847276360538515787675073307545996140586680
9409039 -7.1414511110591131784529507227345610969262291311638790694930667393
74359901464314; 7.498461733909194228616396509663570512888003206562900425148
571916254283480153233 9.476762891025910553896718984259631885515996130393278
289053687511366635168784387; 9.97618115678731426970515029422331603701058582
3224410178838463864917440399374108 21.9936734162634622174628688907960049512

```



```
9850870937463864938968901708055930702034; -4.824183122371212072227360332052
46666913462312409626486941510403920538780421682 -7.327856418070343007524586
384749215033524038596503733266624671149456739680426299]
```

```
sol[1,3]
```

```
-2.622619725152308500938409478465585618829190084067122069098265264160226377
383755
```

To really make use of this, we would want to change `abstol` and `reltol` to be small! Notice that the type for "time" is different than the type for the dependent variables, and this can be used to optimize the algorithm via keeping multiple precisions. We can convert time to be arbitrary precision as well by defining our time span with `BigFloat` variables:

```
prob = ODEProblem(f,big_u0,big.(tspan))
sol = solve(prob)
```

```
retcode: Success
```

```
Interpolation: automatic order switching interpolation
```

```
t: 5-element Array{BigFloat,1}:
```

```
0.0
0.159456998855515426709261384172414149860661134579339478744415384034460600
2778608
0.421191334061567844963495054061999650603920184392629107978780523761969483
5843882
0.721649597012603220345202322675873150865147589172971373295652499255410274
0336311
1.0
```

```
u: 5-element Array{Array{BigFloat,2},1}:
```

```
[0.398481953079453887767158448696136474609375 0.96375568184310078656551468
15697662532329559326171875; 0.914762082592518233070677524665370583534240722
65625 0.361752252114065253607577687944285571575164794921875; 0.626797672347
1335720461183882434852421283721923828125 0.82713264953384046940243479184573
51624965667724609375; 0.809233975242293235297097453440073877573013305664062
5 0.9550891174747071499240291814203374087810516357421875]
[-0.4559188521276427043350085877247006233728638419992953291466322756210757
169659049 -0.23925412414778978844436920353090492924642690660131104956243681
29169838672690548; 0.575680708005470943904124236429961281039199658941607028
3950910984659363432549496 0.22436991924335784294846870316811590371213079678
51045860707345305388087845431292; 0.788644646767375991198262507974780657916
3294579413038440441361556893942313064399 0.78853508225456410399811923493499
21022577928889774360604284286995858235873862633; 1.291542317333483660492401
015135380051822011550541205256990711906753387340738451 2.191008675834062184
617155674946831285941193815996142638905265636076320282061506]
[-2.6226197251523083508194166688466402157254237341358399353104163225474846
14545146 -4.381683442002130509830722414171852300176636045983037838290572305
43392124969933; -0.44396293808566938930908157467556626615557928439822562628
40844660187803509925525 -1.809917286516711726944376045847950456384973605028
47796803777264386758724255558; 2.71515735163651580208972054622201983498293
5344845888198045046179773499455339415 3.74324621797609815922673560446023736
9178260806734825369832371545109872575795371; 1.1896048898492358488241697764
51409082891104405578647439457368192403923338004215 2.9100764381096547092086
5744059638194245928129812745015957028003720307422325701]
[-4.1861253001774564052836881868036268726767540569353393417505547196211362
64397857 -9.273734292428035803630479868738691364749289897784424652157584494
669149055985657; 0.83506057611031867031537086030287221669031524972357865573
6690264905139023716774 -1.2577793278641597685553797309302485849602141767192
69301734507968531831888559079; 7.175255030881521052110458181480522402729295
347251237464430077715176603086044876 12.79300489363077420626645055384150838
```

```

971626173359759852898690780536123583378086; -0.9617338862320839485437094530
914269821358001881403463742289333825616517285939729 -0.05007236416907099910
513736450156282467410357841653594899655130638662799033116438]
[-1.2003236055657169764644456054534449641855468639315850451407140131273725
14347568 -7.141451111059113997182266154449581047477457579675621974429280171
553456617398776; 7.49846173390919326885480614173975981475875889953085442440
8712552013123642771748 9.47676289102590882086820977159584712636418680903130
5131839443038255481030395209; 9.9761811567873142703404802071073179167430256
4139122961468070942756670140877472 21.9936734162634616279914771239299574670
9414431036955078914540415568232028612244; -4.824183122371211641496129891885
93461486994963338509233639636010815287788431156 -7.32785641807034210111066
7480694355134968031044031495522807940000602513449933393]

```

Let's end by showing a more complicated use of types. For small arrays, it's usually faster to do operations on static arrays via the package [StaticArrays.jl](#). The syntax is similar to that of normal arrays, but for these special arrays we utilize the `@SMatrix` macro to indicate we want to create a static array.

```

using StaticArrays
A = @SMatrix [ 1.0  0.0 0.0 -5.0
               4.0 -2.0 4.0 -3.0
              -4.0  0.0 0.0  1.0
               5.0 -2.0 2.0  3.0]
u0 = @SMatrix rand(4,2)
tspan = (0.0,1.0)
f(u,p,t) = A*u
prob = ODEProblem(f,u0,tspan)
sol = solve(prob)

retcode: Success
Interpolation: automatic order switching interpolation
t: 11-element Array{Float64,1}:
 0.0
 0.007163927323932619
 0.04615872311533139
 0.11560313880988551
 0.19730087428935955
 0.29141428170770495
 0.4020236545779373
 0.5329722383938044
 0.6718148245305243
 0.8396162958261761
 1.0
u: 11-element Array{StaticArrays.SArray{Tuple{4,2},Float64,2,8},1}:
 [0.010464809441216527 0.9352435235736136; 0.9228135670166842 0.49975118963
576404; 0.25187072683703304 0.5209977130569305; 0.8953651538109129 0.407290
2249431578]
 [-0.021823451543931616 0.9265627065592806; 0.8975715311192906 0.5244701530
463395; 0.25848166888748125 0.49738837603506664; 0.9051253450889294 0.44980
48962770293]
 [-0.20727560906409667 0.851266068473453; 0.7510290472074644 0.620765380629
9068; 0.31235022214676467 0.3801400524464554; 0.9493515805551066 0.67897061
97675234]
 [-0.5719543793632117 0.5980433770002167; 0.46495118838857386 0.63671966783
85356; 0.4872511896220083 0.23618313730943863; 0.987532015816827 1.06978438
4265273]
 [-1.037698652148758 0.10524211021905411; 0.12567951938435745 0.42305441128
873406; 0.8295910969715014 0.22020152396923182; 0.9573434431961764 1.479158
395177286]

```

```

[-1.581660821715615 -0.7083564551546928; -0.1967274635514562 -0.0698937928
9983344; 1.407318651635367 0.4833618615068218; 0.8086781833186691 1.8435113
091016242]
[-2.1493504019812484 -1.944239057896934; -0.36134490167315436 -0.835074032
8184774; 2.309470758683357 1.278996043289152; 0.4615498439843619 2.05946841
6464487]
[-2.561452699875834 -3.6353405906170813; -0.0650855888861212 -1.6699773509
167564; 3.579756544163601 3.0001673931104116; -0.20371156834908782 1.911856
1388323214]
[-2.447007197602169 -5.365124275834457; 1.0450947549699208 -1.935677165253
8507; 4.9091225070456606 5.733120257263439; -1.1969707674270826 1.147291203
2540945]
[-1.1555184292555893 -6.633300097217812; 3.655133504009725 -0.497832090239
8792; 5.877520119948311 9.885452566028606; -2.68804579751542 -0.75576336517
16178]
[1.6144357555953899 -5.998624122824556; 7.315665827543399 3.54165628381118
67; 5.266445686065155 13.728275763610508; -4.159512101820095 -3.58613086138
64566]

sol[3]

4×2 StaticArrays.SArray{Tuple{4,2},Float64,2,8} with indices SOneTo(4)×SOne
To(2):
-0.207276  0.851266
 0.751029  0.620765
 0.31235   0.38014
 0.949352  0.678971

```

0.3 Conclusion

These are the basic controls in `DifferentialEquations.jl`. All equations are defined via a problem type, and the `solve` command is used with an algorithm choice (or the default) to get a solution. Every solution acts the same, like an array `sol[i]` with `sol.t[i]`, and also like a continuous function `sol(t)` with a nice plot command `plot(sol)`. The Common Solver Options can be used to control the solver for any equation type. Lastly, the types used in the numerical solving are determined by the input types, and this can be used to solve with arbitrary precision and add additional optimizations (this can be used to solve via GPUs for example!). While this was shown on ODEs, these techniques generalize to other types of equations as well.

0.4 Appendix

This tutorial is part of the `SciMLTutorials.jl` repository, found at: <https://github.com/SciMLTutorials/SciMLTutorials.jl>. For more information on doing scientific machine learning (SciML) with open source software, check out <https://sciml.ai/>.

To locally run this tutorial, do the following commands:

```

using SciMLTutorials
SciMLTutorials.weave_file("introduction", "01-ode_introduction.jmd")

```

Computer Information:

```
Julia Version 1.4.2
Commit 44fa15b150* (2020-05-23 18:35 UTC)
Platform Info:
  OS: Linux (x86_64-pc-linux-gnu)
  CPU: Intel(R) Core(TM) i7-9700K CPU @ 3.60GHz
  WORD_SIZE: 64
  LIBM: libopenlibm
  LLVM: libLLVM-8.0.1 (ORCJIT, skylake)
Environment:
  JULIA_LOAD_PATH = /builds/JuliaGPU/DiffEqTutorials.jl:
  JULIA_DEPOT_PATH = /builds/JuliaGPU/DiffEqTutorials.jl/.julia
  JULIA_CUDA_MEMORY_LIMIT = 2147483648
  JULIA_NUM_THREADS = 8
```

Package Information:

```
Status `~/builds/JuliaGPU/DiffEqTutorials.jl/tutorials/introduction/Project.toml`
[6e4b80f9-dd63-53aa-95a3-0cdb28fa8baf] BenchmarkTools 0.5.0
[0c46a032-eb83-5123-abaf-570d42b7fbaf] DifferentialEquations 6.15.0
[65888b18-ceab-5e60-b2b9-181511a3b968] ParameterizedFunctions 5.4.0
[91a5bcdd-55d7-5caf-9e0b-520d859cae80] Plots 1.5.6
[90137ffa-7385-5640-81b9-e52037218182] StaticArrays 0.12.4
[c3572dad-4567-51f8-b174-8c6c989267f4] Sundials 4.2.5
[37e2e46d-f89d-539d-b4ee-838fcccc9c8e] LinearAlgebra
```